

# Probability Basics and Linear Classification

Machine Learning – Basic Methods

Prof. Gerhard Neumann

Autonomous Learning Robots

KIT, Institut für Anthropomatik und Robotik

# Learning Outcomes

- Understand probabilistic models and maximum likelihood
- Understand the classification problem
- What is a linear classifier?
- What is the loss function of linear classification?
- What is gradient descent ?

# Today's Agenda!

## **Basics: Probability Theory**

- Probabilistic Models
- Expectations and Monte Carlo Methods
- Maximum Likelihood

## **Basics: Gradient Descent**

## **Classification:**

- Generative vs. discriminative classification
- Linear Classification
- Logistic Regression

Many slides are based on slides from Shenlon Wang, Yingyu Jiang, Michail Michailidis and Patrick Maiden

# Basics: Probability Theory

- *“Probability theory is nothing but common sense reduced to calculation”, Pierre Laplace, 1812*
- We will keep our discussion relatively informal and pick the things we need from probability theory

# Notation

- A **random variable**  $X$  represents uncertain states or outcomes of the world
- We will write  $p(x)$  to mean the probability that  $X$  takes the value  $x$
- The sample space is the space of all possible outcomes
  - Might be discrete, continuous or mixed
- $p(x)$  is the **probability mass** (density) function
  - Assigns a number to each point of the sample space
  - Non-negative, sums (integrates) to 1
  - Intuitively: How often does  $x$  occur? How much do we believe in  $x$ ?

# Distributions

- **Joint distribution**

$$p(x, y)$$

- Probability that  $X=x$  and  $Y=y$

- **Conditional distribution**

$$p(x|y)$$

- Probability that  $X=x$  given  $Y=y$

Conditional Distributions

$P(W T)$	$P(W T = \text{hot})$		
		W	P
		sun	0.8
		rain	0.2
	$P(W T = \text{cold})$		
		W	P
		sun	0.4
		rain	0.6

Joint Distribution

$P(T, W)$		
T	W	P
hot	sun	0.4
hot	rain	0.1
cold	sun	0.2
cold	rain	0.3

# Rules of Probability

- **Sum rule (marginalization / integrating out):**

$$p(x) = \sum_y p(x, y)$$

$$p(x_1) = \sum_{x_2} \sum_{x_3} \cdots \sum_{x_D} p(x_1, \dots, x_D)$$

- **Note:** For continuous distributions, the sums will be replaced by integrals

$P(T, W)$			$P(T)$	
T	W	P		
hot	sun	0.4	hot	0.5
hot	rain	0.1	cold	0.5
cold	sun	0.2	$P(W)$	
cold	rain	0.3	sun	0.6
			rain	0.4

$P(t) = \sum_w P(t, w)$

$P(w) = \sum_t P(t, w)$

# Rules of Probability

- Chain / product rule

$$p(x, y) = p(x|y)p(y)$$

$$p(x_1, \dots, x_D) = p(x_1)p(x_2|x_1) \dots p(x_D|x_1, \dots, x_{D-1})$$

$P(W)$		$P(D W)$			$P(D, W)$		
W	P	D	W	P	D	W	P
sun	0.8	wet	sun	0.1	wet	sun	0.08
rain	0.2	dry	sun	0.9	dry	sun	0.72
		wet	rain	0.7	wet	rain	0.14
		dry	rain	0.3	dry	rain	0.06



# Bayes Rule

**Bayes rule** is one of the most important equations in probability theory and in machine learning

$$p(x|y) = \frac{p(y|x)p(x)}{p(y)} = \frac{p(y|x)p(x)}{\sum_{x'} p(y|x')p(x')}$$

- Way of “reversing” the conditional probabilities
- Often one conditional is tricky but the other one is simple
- One of the **most important equations** for ML!



# Expectations

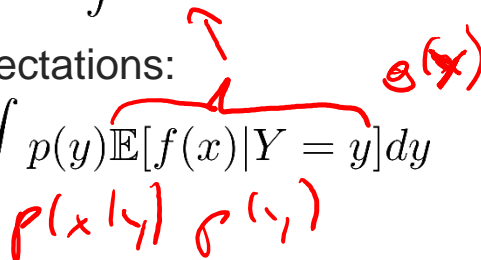
The expectation of a function  $f(x)$  with respect to a distribution  $p(x)$  is given by

$$\mathbb{E}_p[f(x)] = \int p(x)f(x)dx$$

A conditional expectation is given by

$$\mathbb{E}_p[f(x)|Y = y] = \int p(x|y)f(x)dx$$

Chain rule for expectations:

$$\mathbb{E}_p[f(x)] = \int p(y)\mathbb{E}[f(x)|Y = y]dy$$


# Monte-carlo estimation

Expectations can always be **approximated by samples**:

$$\mathbb{E}_p[f(x)] = \int p(x)f(x)dx \approx \frac{1}{N} \sum_{x_i \sim p(x)} f(x_i)$$

- Necessary if no analytical solution exists to compute the integral (typical case)

# Moments

## Moments are expectations:

- 1<sup>st</sup> moment (mean):  $\boldsymbol{\mu} = \mathbb{E}_p[\boldsymbol{x}]$
- 2<sup>nd</sup> moment:  $\boldsymbol{M}_2 = \mathbb{E}_p[\boldsymbol{x}\boldsymbol{x}^T]$

## Central moments are always computed relatively to the mean:

- 2<sup>nd</sup> central moment (covariance):
$$\boldsymbol{\Sigma} = \mathbb{E}_p[(\boldsymbol{x} - \boldsymbol{\mu})(\boldsymbol{x} - \boldsymbol{\mu})^T]$$
- Captures variability (diagonal entries) and correlation (off-diagonal)

# Distributions

## Bernoulli Distribution:

- Binary random variable  $X \in \{0, 1\}$
- One parameter  $p(X = 1) = \mu$
- Probability distribution  $p(x) = \mu^x (1 - \mu)^{(1-x)}$
- Think of it as tossing a coin

Depending on  $x$ , selects either  $\mu$  or  $1-\mu$  as probability

# Distributions

## Multinomial / Categorical Distribution:

- K different events:  $C \in \{1, \dots, K\}$
- Directly specifies probabilities:  $p(C = k) = \mu_k, \quad \mu_k \geq 0, \quad \sum_{k=1}^K \mu_k = 1$
- Or written with 1-hot-encoding (without an “if” clause)

$$p(c) = \prod_{k=1}^K \mu_k^{h_{c,k}}$$

$h_c = [1, 0, 0, \dots, 0]$ ,  $h_c = [0, 0, 1, 0, \dots]$   
Depending on the class label of  $x$ , selects the correct  $\mu_k$


- where  $h_x$  is the K-dimensional 1-hot encoding vector, which is one for the dimension  $c = k$  and 0 elsewhere.  $h_{x,k}$  is the k-th element of this vector.
- Think of it as tossing a die

# Distributions

## Gaussian Distribution

- Continuous RV:  $X \in \mathbb{R}$
- Distribution is completely specified by mean  $\mu$  and variance  $\sigma^2$

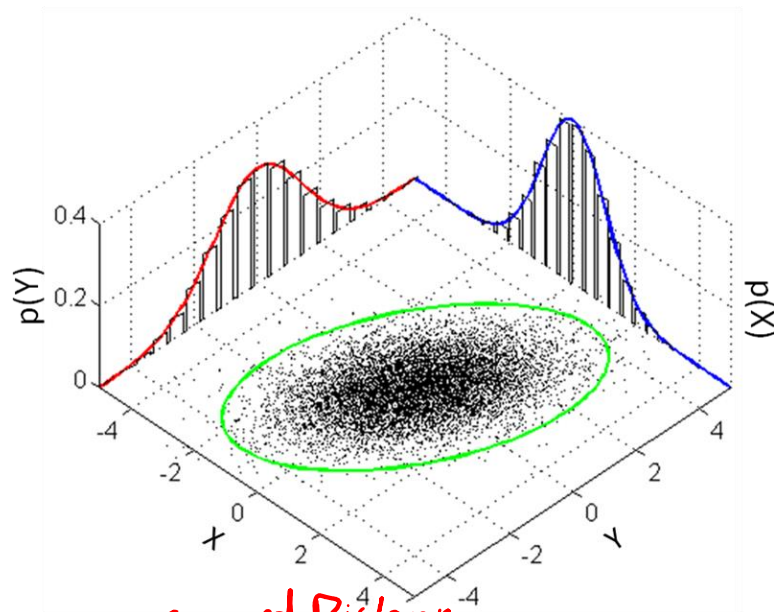
$$p(x) = \mathcal{N}(x|\mu, \sigma) = \frac{1}{\sqrt{2\pi\sigma^2}} \exp \left\{ -\frac{(x - \mu)^2}{2\sigma^2} \right\}$$



# Distributions

## Multivariate Gaussian Distribution

- Continuous RV:  $X \in \mathbb{R}^d$
- Distribution is completely specified by mean vector  $\mu$  and covariance matrix  $\Sigma$



$$p(x) = \mathcal{N}(x|\mu, \Sigma) = \frac{1}{\sqrt{|2\pi\Sigma|}} \exp \left\{ -\frac{((x - \mu)^T \Sigma^{-1} (x - \mu))}{2} \right\}$$

*Squared Distance*

*Normalization*



# Distributions

## **Important Properties of Gaussians:**

- All marginals of a Gaussian are again Gaussian
- Every conditional is Gaussian
- The product of 2 Gaussians is again Gaussian
- Even the sum of 2 Gaussian RVs is again Gaussian

# Maximum Likelihood Estimation (MLE)

- Given: the training data  $D = \{(x_i, y_i)\}_{i=1 \dots N}$  **identically independently distributed (iid)** from the data distribution  $p_{data}$
- Let  $p_{\theta}(x, y)$  be a family of distributions parametrized by  $\theta \in \Theta$
- We want to find  $\theta$  such that  $p$  fits the data well

**Fitness of  $\theta$  for one single data point:**

$$\text{lik}(\theta; x_i, y_i) = p_{\theta}(x_i, y_i)$$

**Fitness of  $\theta$  for whole dataset (iid. assumption):**

$$\text{lik}(\theta; D) = \prod_i p_{\theta}(x_i, y_i)$$

# Maximum Likelihood Estimation (MLE)

**Log-likelihood is easier to optimize:**

$$\text{loglik}(\boldsymbol{\theta}; D) = \sum_i \log p_{\boldsymbol{\theta}}(x_i, y_i)$$

- Log is monotonous -> same optimum
- Sums are “nicer” to optimize than products
- Log cancels exponential form (most distributions are in the exponential family)

**The MLE solution is given by:**

$$\boldsymbol{\theta}_{\text{ML}} = \text{argmax}_{\boldsymbol{\theta}} \text{loglik}(\boldsymbol{\theta}; D)$$

# Example: Gaussian

**Gaussian density function:**

$$\text{loglik}(\boldsymbol{\theta}; D) = -N \log \sqrt{2\pi\sigma^2} - \sum_i \frac{(x_i - \mu)^2}{2\sigma^2}$$

**MLE solution for  $\mu$ :**

# MLE: conditional log-likelihood

- Given the training data  $D = \{(x_i, y_i)\}_{i=1 \dots N}$  iid. from the data distribution  $p_{data}$
- Let  $p_{\theta}(y|x)$  be a family of distributions parametrized by  $\theta \in \Theta$
- We only care about distribution of  $y$ , not of  $x$
- Typical case in supervised learning

## Log-likelihood:

$$\text{loglik}(\theta; D) = \sum_i \log p_{\theta}(y_i|x_i)$$

# Example: Linear Gaussian model

We consider the following conditional Gaussian model:

$$p_{\boldsymbol{\theta}}(y|\mathbf{x}) = \mathcal{N}(y|\mathbf{w}^T \tilde{\mathbf{x}}, \sigma^2), \quad \boldsymbol{\theta} = \{\mathbf{w}, \sigma^2\}$$

**Log-likelihood:**

$$\text{loglik}(\boldsymbol{\theta}; D) = -\log \sqrt{2\pi\sigma^2} - \sum_i \frac{(y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i)^2}{2\sigma^2}$$

- For obtaining  $\mathbf{w}$ , only the squared errors matter, i.e.

$$\text{loglik}(\boldsymbol{\theta}; D) = \text{const}_1 - \text{const}_2 \sum_i (y_i - \mathbf{w}^T \tilde{\mathbf{x}}_i)^2$$

- Hence, the **MLE solution is equivalent to the least squares solution!**
- **But:** we can also obtain the variance!

# Takeaway messages

## What have we learned so far?

- Basic rules of probabilities ... nothing new so far
- Expectations can be evaluated by samples
- How to compute the ML estimator
- Maximum likelihood is equivalent to minimizing the squared loss for:
  - Conditional Gaussian models
  - With constant noise



# Today's Agenda!

## **Basics: Probability Theory**

- Probabilistic Models
- Expectations and Monte Carlo Methods
- Maximum Likelihood

## **Linear Classification:**

- Linear Classifiers
- Logistic Regression

## **Basics: Gradient Descent**

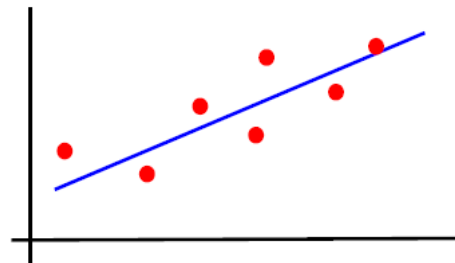


# Supervised Learning

Training data includes targets

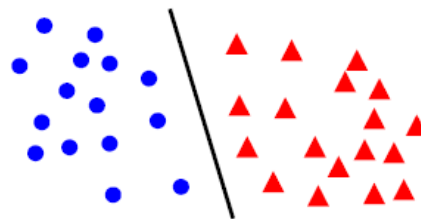
- **Regression:**

- Learn continuous function
- Example: line

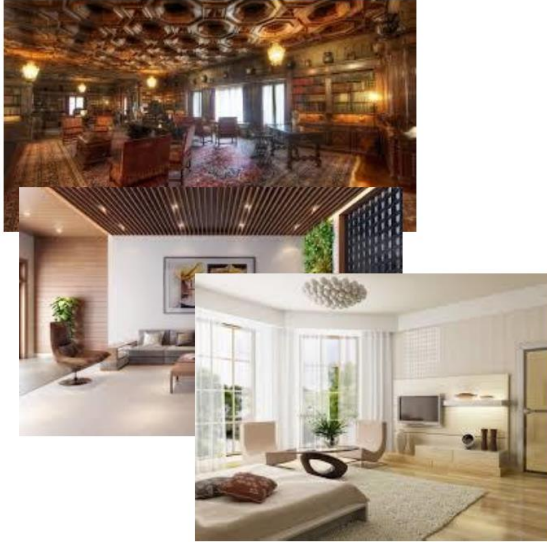


- **Classification:**

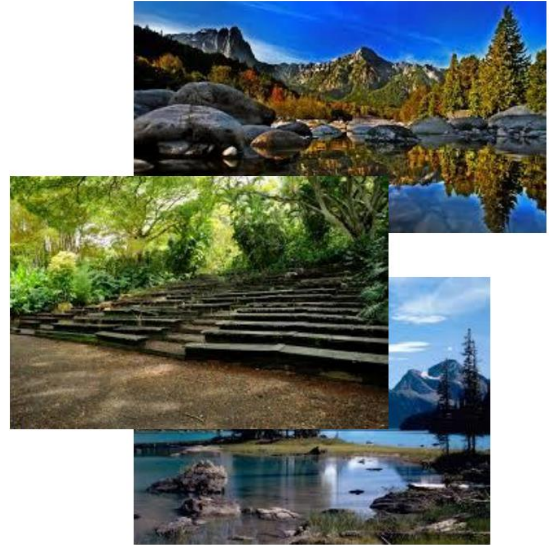
- Learn class labels
- Example: Digit recognition



# Example 1: Image classification



Indoor



outdoor

## Example 2: Spam Classification

	#"\$"	#"Mr."	#"sale"	...	Spam?
Email 1	2	1	1		Yes
Email 2	0	1	0		No
Email 3	1	1	1		Yes
...					
Email n	0	0	0		No
New email	0	0	1		??

# Definition

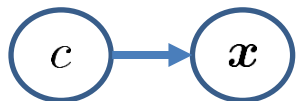
Given the dataset  $\mathcal{D} = \{(\mathbf{x}_i, c_i)\}_{i=1\dots N}$ , where  $\mathbf{x}_i \in \mathbb{R}^d$  are the input samples and  $c \in \{1 \dots K\}$  are the class labels, we want to learn a classifier  $f(\mathbf{x})$  that predicts the class label for unseen samples.

- $K = 2$ : Binary classification
- $K > 2$ : Multi-class classification

In difference to regression, the **output is now discrete!**

# Generative vs. discriminative modelling

## Generative Models:




- Assume some functional form for class prior  $p(c)$  and class densities  $p(x|c)$
- Learn prior and densities from data
  - This is a “generative” model, as we can create new datapoints  $x$  using  $p(x|c)$
- Predict class label by **computing posterior**  $p(c|x) = \frac{p(x|c)p(c)}{p(x)}$

## Learn full **joint distribution** of the data (typically very **hard**)

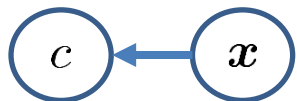
- Our modelling assumptions, e.g. that  $p(x|c)$  is Gaussian, might introduce big errors

## Generative



# Generative vs. discriminative modelling

## Discriminative Models:



- Directly assume some functional form for  $p(c|x)$  (or any other predictor  $f(x)$  that returns the class label).
- This is a ‘discriminative’ model of the data!
- Estimate parameters of  $p(c|x)$  directly from training data

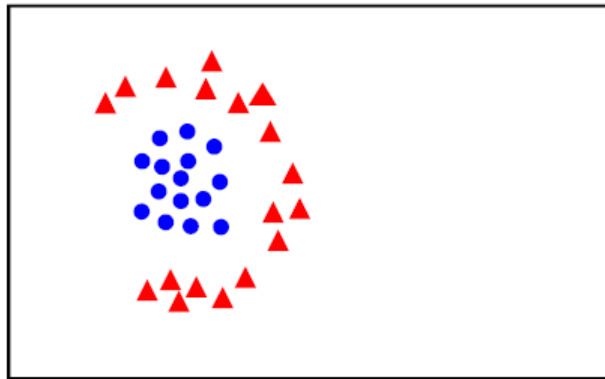
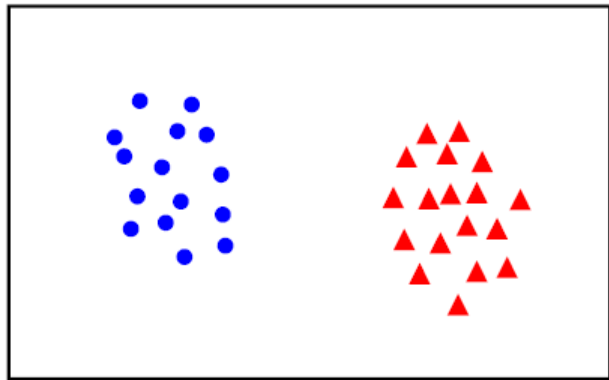
## Modelling needs to **consider only points on the border**

- Typically much simpler than generative modelling
- We therefore concentrate on discriminative models

## Discriminative



# (Discriminative) Binary Classification



Given the training data  $(\mathbf{x}_i, y_i)$ ,  $i = 1 \dots N$ , with  $\mathbf{x}_i \in \mathbb{R}^d$  and  $y_i \in \{0, 1\}$ , learn a classifier  $f(\mathbf{x})$  such that:

$$f(\mathbf{x}_i) = \begin{cases} > 0, & \text{if } y_i = 1 \\ < 0, & \text{if } y_i = 0 \end{cases}$$

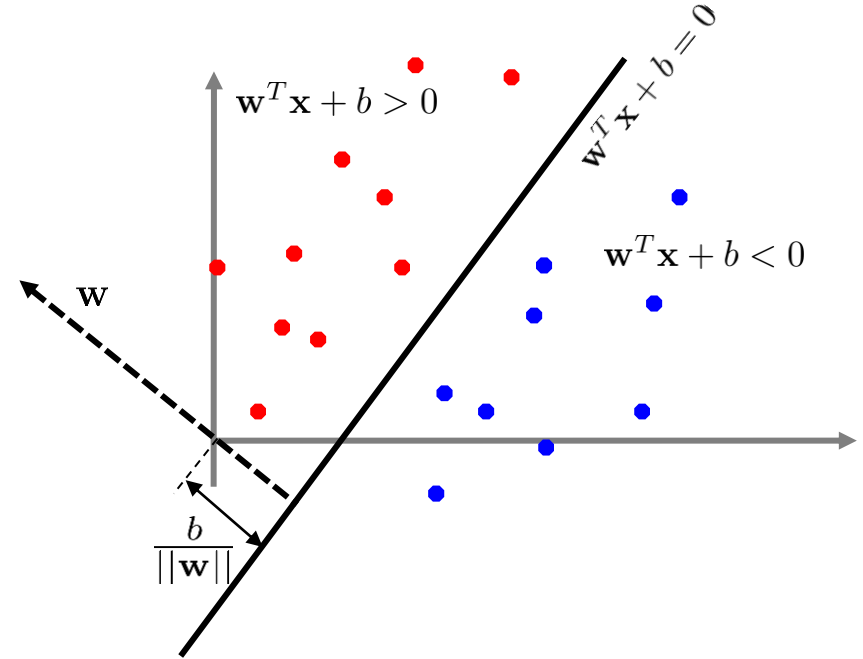
# Linear Classifiers

**A linear classifier is given in the form:**

$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

**In 2D, the classifier is a line**

- $\mathbf{w}$  is the normal to the line
- $b$  is the bias





# Linear Discriminators

A linear discriminator is given in the form:


$$f(\mathbf{x}) = \mathbf{w}^T \mathbf{x} + b$$

In 2D, the discriminator is a **line**

- $\mathbf{w}$  is the normal to the line
- $b$  is the bias

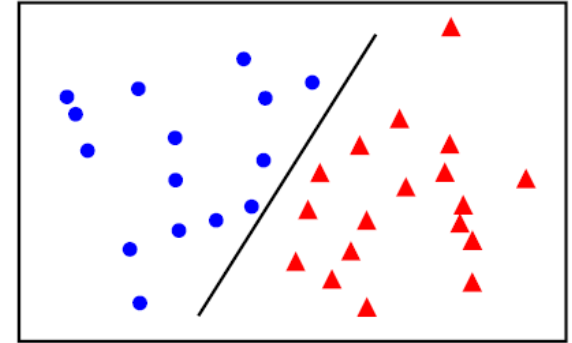
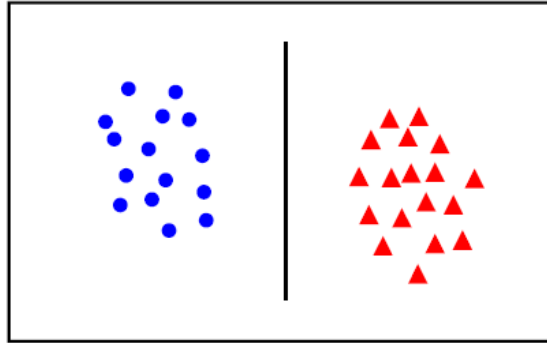
In 3D, it's a **plane**

In N-D, it's a **hyper-plane**

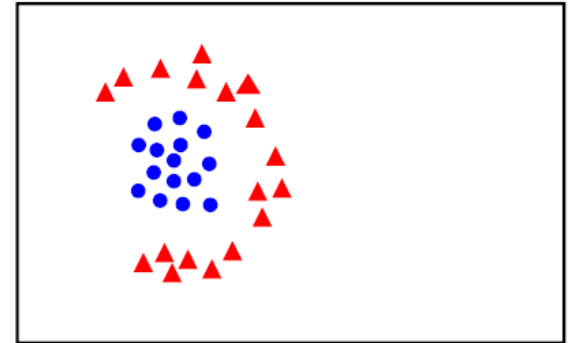
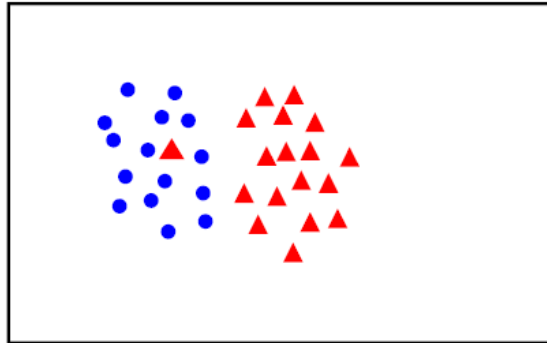


# Linear Separability

**Linear Separable**



**Non-Linear Separable**



# Linear Classification: 0-1 loss (1<sup>st</sup> attempt)

**Prediction:**  $y = \text{step}(f(\mathbf{x})) = \text{step}(\mathbf{w}^T \mathbf{x} + b)$

- Predict class 1 for  $f(\mathbf{x}) > 0$
- else predict class 0

**Optimization:** Find  $\mathbf{w}$  such that

$$L_0(\mathbf{w}) = \sum_i \mathbb{I}(\text{step}(\mathbf{w}^T \mathbf{x} + b) \neq y_i)$$

- where  $\mathbb{I}$  returns 1 if the argument is true
- ... counts the number of misclassifications

✗ Very **difficult to optimize!!!** (NP-hard)

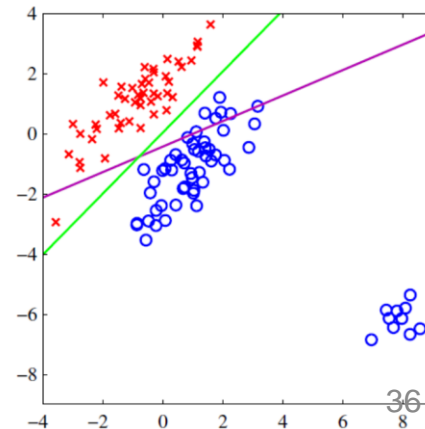
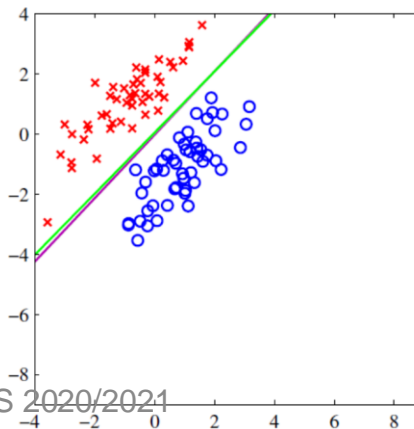
# Linear Classification: regression loss (2<sup>nd</sup> attempt)

**We can use same loss as in regression**

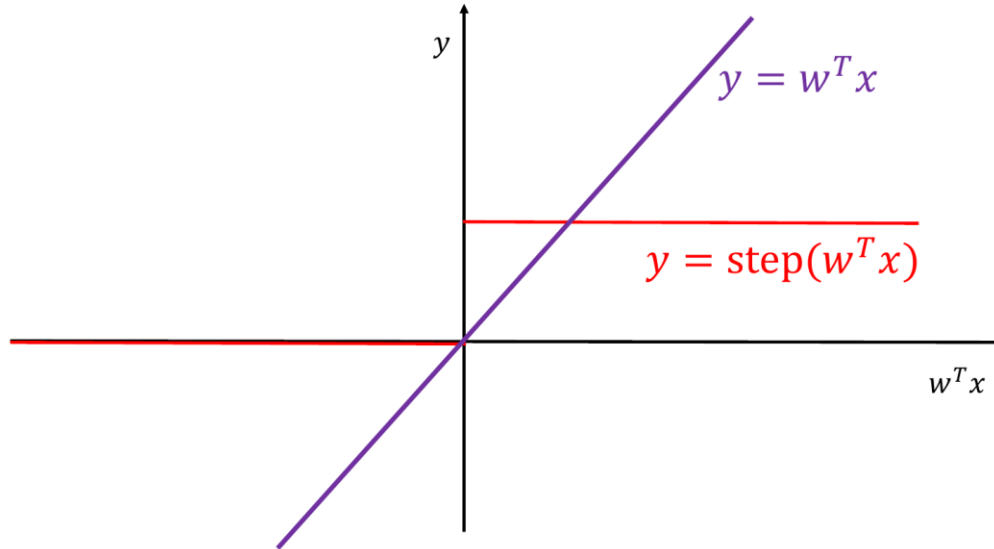
$$L_{\text{reg}}(\mathbf{w}) = \sum_i (f(\mathbf{x}_i) - y_i)^2$$

- Minimize the squared error: Easy!
- However: we ignored the fact that  $y_i$  is restricted to  $\{0,1\}$

× **Not robust** to outliers



# Compare the two



- The output of a linear function is unbounded!
- However, useful output values are only 0 or 1

# Logistic sigmoid function

## Sigmoid function:

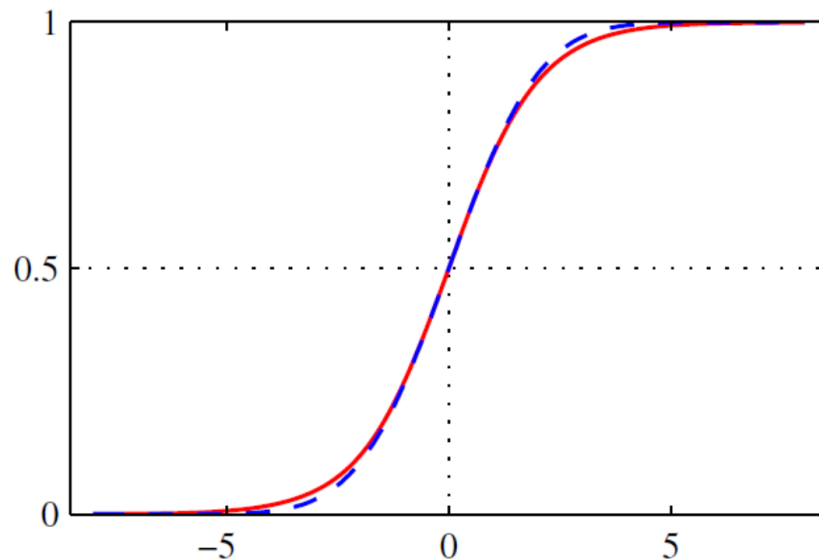
$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

- Output is bounded between 0 and 1
- Smooth

## For linear classification:

- Squash the output of the linear function
- Minimize the loss

$$L(\mathbf{w}) = \sum_i (\sigma(f(\mathbf{x}_i)) - y_i)^2 = \sum_i (\sigma(\mathbf{w}^T \mathbf{x} + b) - y_i)^2$$



# Better: Probabilistic View

## Define conditional probability distribution of the class label

$$p(c = 1|\mathbf{x}) = \sigma(\mathbf{w}^T \mathbf{x} + b), \quad p(c = 0|\mathbf{x}) = 1 - \sigma(\mathbf{w}^T \mathbf{x} + b)$$

- This is now a **conditional Bernoulli** distribution. I.e. the **outcome of the event  $c$  depends on  $\mathbf{x}$**
- We can use the same “**exponential trick**” to select the correct probability depending on the value of  $c$ , i.e.

$$p(c|\mathbf{x}) = p(c = 1|\mathbf{x})^c p(c = 0|\mathbf{x})^{1-c} = \sigma(\mathbf{w}^T \mathbf{x} + b)^c (1 - \sigma(\mathbf{w}^T \mathbf{x} + b))^{1-c}$$

# Log-Likelihood

**We can now directly optimize the conditional Bernoulli log-likelihood**

$$\begin{aligned}\text{loglik}(\tilde{\mathbf{w}}, D) &= \sum_i \log p(c_i | \mathbf{x}_i) = \sum_i \log (p(c = 1 | \mathbf{x}_i)^{c_i} p(c = 0 | \mathbf{x}_i)^{1-c_i}) \\ &= \sum_i c_i \log p(c = 1 | \mathbf{x}_i) + (1 - c_i) \log p(c = 0 | \mathbf{x}_i) \\ &= \sum_i c_i \log \sigma(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i) + (1 - c_i) \log (1 - \sigma(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i))\end{aligned}$$

- Negative likelihood is also often referred to as **cross-entropy loss**



# Logistic Regression

Optimizing the log-likelihood of a sigmoid is called **logistic regression**

$$\operatorname{argmax}_{\tilde{\mathbf{w}}} \operatorname{loglik}(\tilde{\mathbf{w}}, D) = \operatorname{argmax}_{\tilde{\mathbf{w}}} \sum_i c_i \log \sigma(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i) + (1 - c_i) \log (1 - \sigma(\tilde{\mathbf{w}}^T \tilde{\mathbf{x}}_i))$$

- ... even though we solve a classification problem
- One can show that the function is still convex (only one maximum exists)
- However, there is **no closed form solution** as in linear regression

How can we find the maximum? -> **Gradient Descent!**

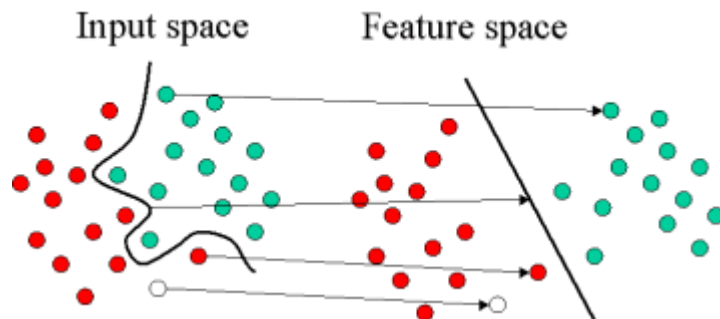
# Generalized logistic models

We can fit a **linear discriminator** in a **non-linear feature space**

- Similar to generalized linear regression models

$$\operatorname{argmax}_{\mathbf{w}} \operatorname{loglik}(\mathbf{w}, D) = \operatorname{argmax}_{\mathbf{w}} \sum_i c_i \log \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)) + (1 - c_i) \log (1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)))$$

- Problems that are not linear separable in input space can be linear separable in feature space



# Regularization

Similar as in linear regression, we can again add a **regularization penalty**

$$L(\tilde{\mathbf{w}}, D) = \text{loglik}(\tilde{\mathbf{w}}, D) - \lambda \text{penalty}(\tilde{\mathbf{w}})$$

**Most common: L2 regularization loss**

$$\text{penalty}(\tilde{\mathbf{w}}) = ||\tilde{\mathbf{w}}||^2$$

- L is still convex for most penalty terms

# Today's Agenda!

## **Basics: Probability Theory**

- Probabilistic Models
- Expectations and Monte Carlo Methods
- Maximum Likelihood

## **Linear Classification:**

- Linear Classifiers
- Logistic Regression

## **Basics: Gradient Descent**

# Optimization

For most ML algorithms, we want to find the best model to fit the data.

**Two examples we already know:**

- **Least squares solution:**

$$\operatorname{argmin}_{\boldsymbol{w}} \operatorname{SSE}(\boldsymbol{w}, D)$$

- **Maximum likelihood solution:**

$$\operatorname{argmax}_{\boldsymbol{w}} \operatorname{loglik}(\boldsymbol{w}, D)$$

# Optimization

For most ML algorithms, we want to find the best model to fit the data.

**Two examples we already know:**

- **Least squares solution:**

$$\operatorname{argmin}_{\boldsymbol{w}} \operatorname{SSE}(\boldsymbol{w}, D) + \lambda \operatorname{penalty}(\boldsymbol{w})$$

- **Maximum likelihood solution:**

$$\operatorname{argmax}_{\boldsymbol{w}} \operatorname{loglik}(\boldsymbol{w}, D) - \lambda \operatorname{penalty}(\boldsymbol{w})$$

**... plus regularization penalty**

Note that:

$$\operatorname{argmin}_{\boldsymbol{x}} f(\boldsymbol{x}) = \operatorname{argmax}_{\boldsymbol{x}} -f(\boldsymbol{x})$$

Hence, the role of the penalty is the same

# Optimization

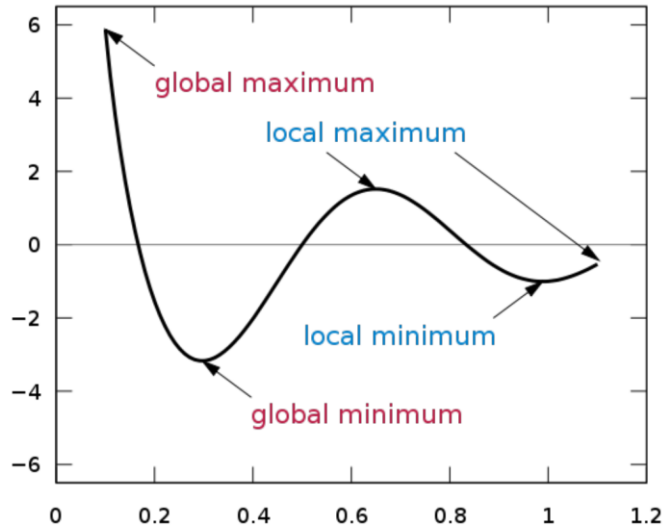
General form of optimization for ML: loss + penalty

$$\arg \min_{\text{parameters } \boldsymbol{\theta}} \sum_{i=1}^N l(\mathbf{x}_i, \boldsymbol{\theta}) + \lambda \text{ penalty}(\boldsymbol{\theta})$$

- Summed sample-loss plus regularization penalty

How to do that? **Optimization**

# When can we do that?



- The global minimum/maximum can only be found for convex functions!
- For non-convex functions we are limited to finding a local minimum / maximum

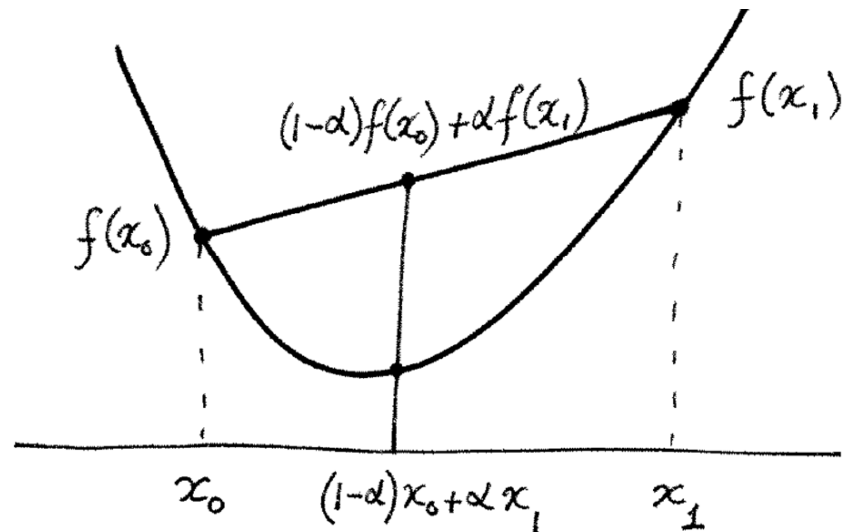


# Convex functions

A convex function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  satisfies for any  $\mathbf{x}_0, \mathbf{x}_1 \in \mathbb{R}^d$

$$f((1 - \alpha)\mathbf{x}_0 + \alpha\mathbf{x}_1) \leq (1 - \alpha)f(\mathbf{x}_0) + \alpha f(\mathbf{x}_1), \quad \alpha \in [0, 1]$$

- Line joining  $(\mathbf{x}_0, f(\mathbf{x}_0))$  and  $(\mathbf{x}_1, f(\mathbf{x}_1))$  is always above the function value
- There is only one minimum!



## Example: Linear Regression Objective

$$L_{\text{ridge}} = (\mathbf{y} - \Phi \mathbf{w})^T (\mathbf{y} - \Phi \mathbf{w}) + \lambda \mathbf{w}^T \mathbf{w}$$

- Convex
- Quadratic function in  $\mathbf{w}$
- Minimum can be obtained analytically
- One of the very rare cases!

In most other cases, we have to resort to incremental methods: **Gradient descent**

# Gradient Descent

- Is good for finding **global minima** if function is **convex**
- Is good for finding **local minima** if function is **non-convex**
- Has many applications in ML:
  - Logistic Regression
  - Linear Regression (for large input dimensions)
  - Neural Networks
  - Mixture Models
  - ...

# Gradient Descent

Start at some point, follow the gradient towards (a) minimum

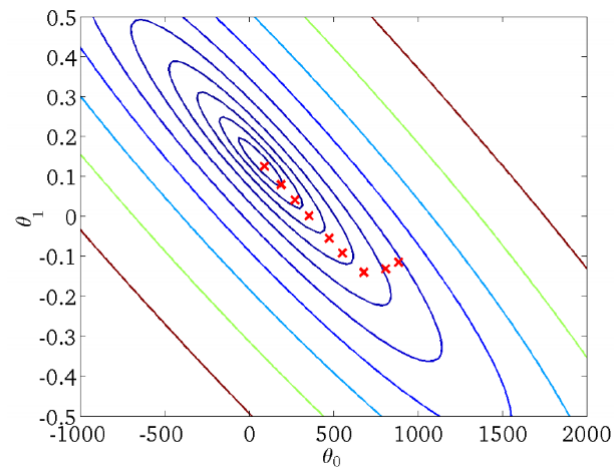
$\mathbf{x}_0 \leftarrow \text{init}, t = 0$

**while** termination condition does not hold **do**

$\mathbf{x}_{t+1} = \mathbf{x}_t - \eta \nabla f(\mathbf{x}_t), \quad t = t + 1$

**end while**

- $\eta \dots$  learning rate or step size
- Gradient always points in the direction of **steepest ascen.**



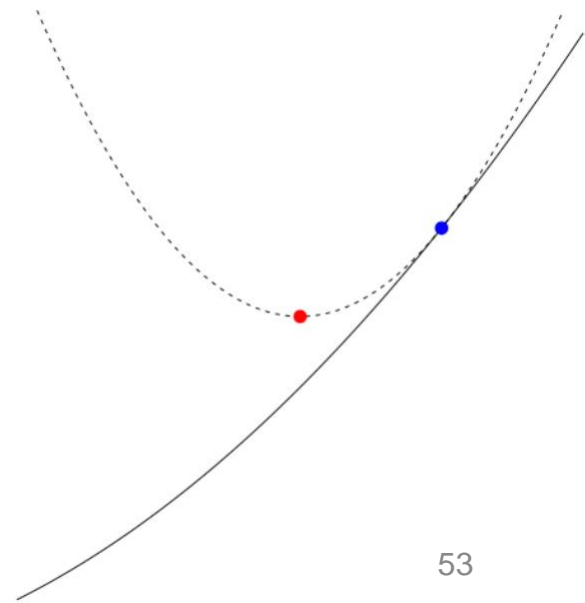
# Gradient Descent Interpretation

Approximate the function as quadratic function:

$$\hat{f}(\mathbf{x}) = \underbrace{f(\mathbf{x}_t) + \nabla f(\mathbf{x}_t)^T (\mathbf{x} - \mathbf{x}_t)}_{\text{linear approximation}} + \underbrace{\frac{1}{2\eta} \|\mathbf{x}_t - \mathbf{x}\|^2}_{\text{proximity of } \mathbf{x}_t} \approx f(\mathbf{x})$$

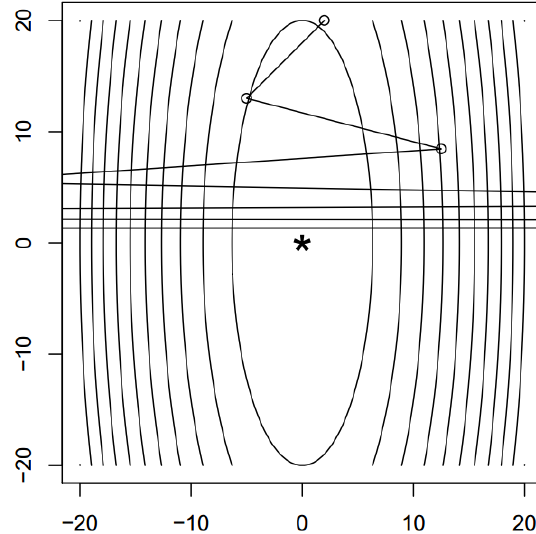
- Finding the minimum of  $\hat{f}(\mathbf{x})$  yields the gradient descent rule

$$\mathbf{x}_{t+1} = \arg \min_{\mathbf{x}} \hat{f}(\mathbf{x}), \quad \mathbf{x}_{t+1} = \mathbf{x}_t - \eta \nabla f(\mathbf{x}_t)$$

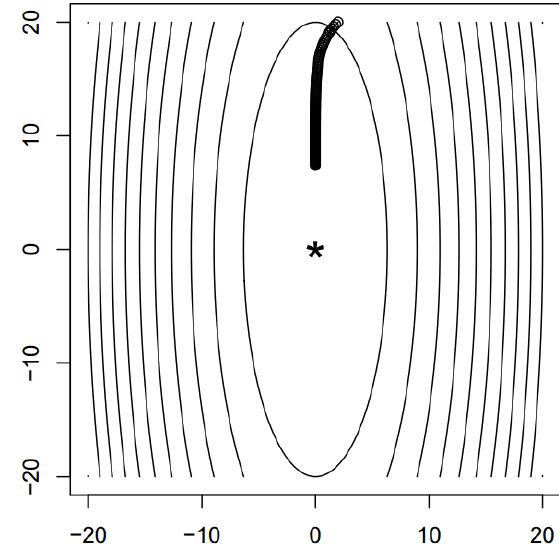


# Choosing the step-size

$\eta_t = t$ , it is too big



too small  $\eta_t$ , after 100 iterations



# How to terminate

## **When change in iterates is small**

- When gradient is small
- When change in function value is small

## **Or after a fixed time step or budget**

# Stochastic Gradient Descent

- Usually we are minimizing the empirical loss (**batch gradient descent**)

$$\frac{1}{n} \sum_i l(\mathbf{x}_i; \boldsymbol{\theta}) \qquad \boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \frac{\eta}{n} \sum_i \nabla_{\boldsymbol{\theta}} l(\mathbf{x}_i; \boldsymbol{\theta}_t)$$

- We do this to approximate the expected loss

$$\mathbb{E}_{\mathbf{x}} [l(\mathbf{x}; \boldsymbol{\theta})] \qquad \boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta \mathbb{E}_{\mathbf{x}} [\nabla_{\boldsymbol{\theta}} l(\mathbf{x}; \boldsymbol{\theta}_t)]$$

- Use a rougher, cheaper approximation: **stochastic gradient descent**

$$l(\mathbf{x}_i; \boldsymbol{\theta}) \qquad \boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \eta \nabla_{\boldsymbol{\theta}} l(\mathbf{x}_i; \boldsymbol{\theta}_t)$$

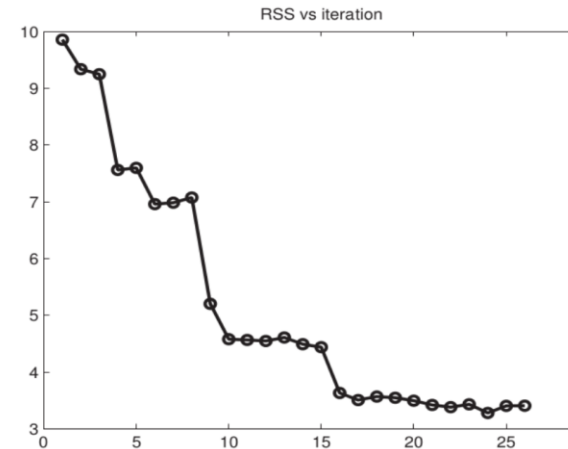
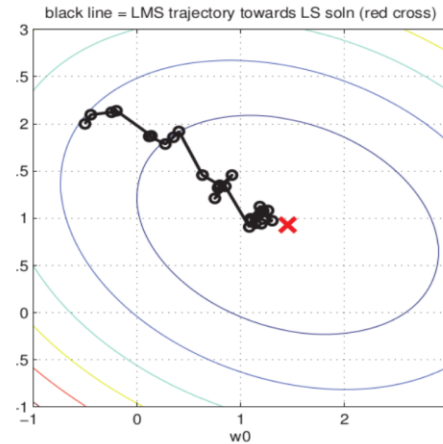
- for random sample  $i$



# Stochastic Gradient Descent (SGD)

## Use only one sample to compute the update

- Does NOT always “descent”
- Iterations are much cheaper
- Requires more iterations
- ... and smaller step sizes

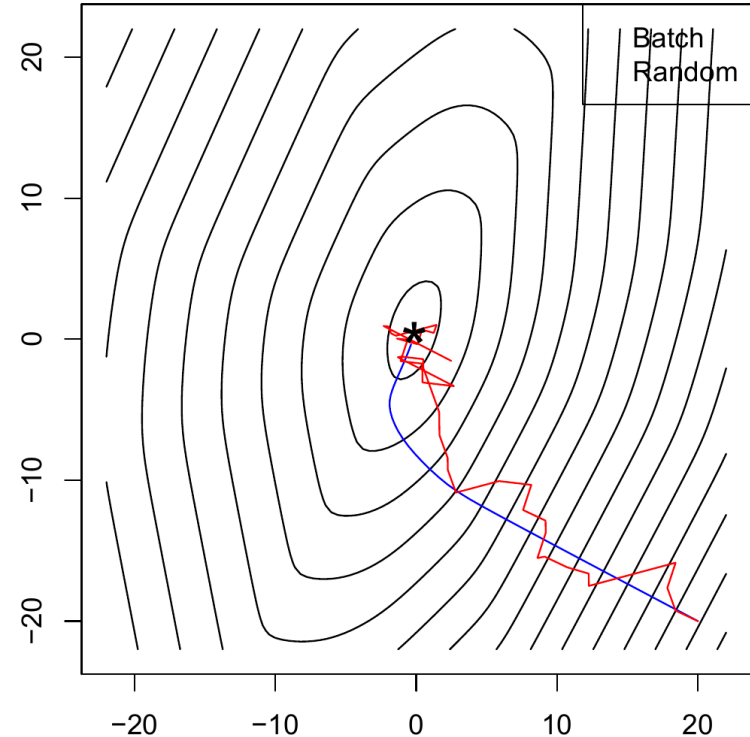


# Stochastic vs. Batch Gradients

- **Blue:** Batch Gradients
- **Red:** Stochastic Gradients

## Rule of thumb:

- Stochastic methods work well far away from optimum
- But struggle to find the exact optimum



# Step-sizes

Standard in SGD is to use diminishing step sizes, e.g.,  $\eta_t = \frac{1}{t}$

- Asymptotically approach the optimum
- instead of “wiggling” around optimum

In general, it can be shown that SGD **converges to the optimum** for strictly convex functions if (**stochastic approximation theory**)

$$\sum_t \eta_t = \infty \quad \text{and} \quad \sum_t \eta_t^2 < \infty$$

# Stochastic vs. Batch Gradients

## Why are stochastic gradients often better than batch?

- Typically, our data-set will contain redundancy
- Hence, some computation in the batch gradients are redundant
  - compute the gradients for similar samples
  - using the same parameter vector
- This does not happen if we update immediately after one sample

As a consequence, **SGD requires less computation** (in most cases)

# Mini-Batches

Take **subset of samples**  $I_t \subset \{1, \dots, n\}$ ,  $|I_t| = b$ ,  $b \ll n$  to approximate real gradient:

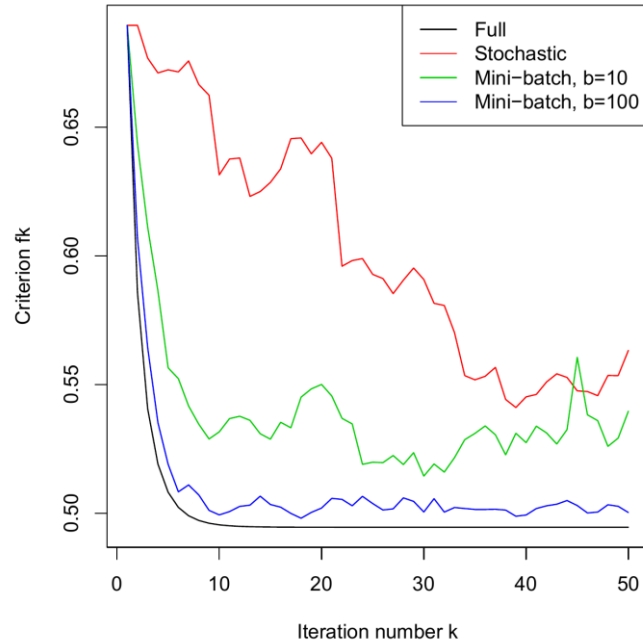
$$\frac{1}{b} \sum_{i \in I_t} l(\mathbf{x}_i; \boldsymbol{\theta})$$

$$\boldsymbol{\theta}_{t+1} = \boldsymbol{\theta}_t - \frac{\eta}{b} \sum_{i \in I_t} \nabla_{\boldsymbol{\theta}} l(\mathbf{x}_i; \boldsymbol{\theta}_t)$$

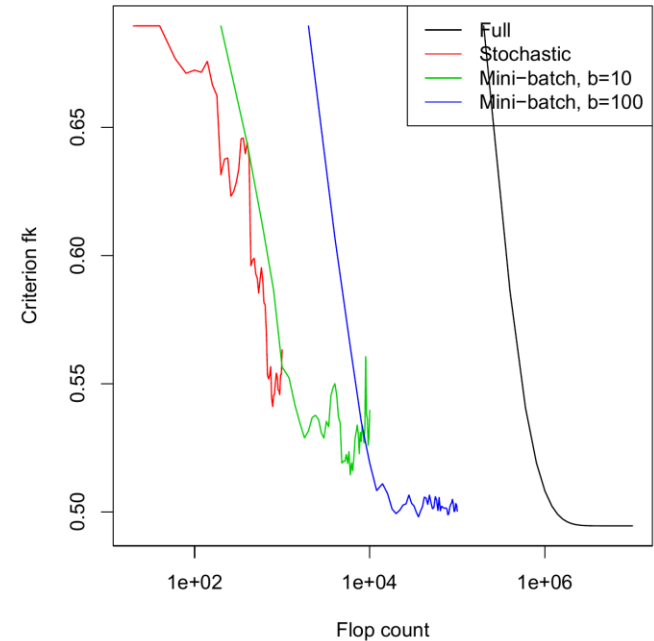
- Intermediate version of stochastic and batch gradient descent
- Less noisy estimates
- Achieves “descent” more often
- Preferable for GPU implementations

# Example

10000 samples, loglikelihood logistic regression:



per iteration



per flop

# Gradient Descent for Logistic Regression

## Properties of the sigmoid function:

- Bounded:  $\sigma(a) = \frac{1}{1 + \exp(-a)} \in (0, 1)$
- Symmetric:  $1 - \sigma(a) = \frac{\exp(-a)}{1 + \exp(-a)} = \frac{1}{1 + \exp(a)} = \sigma(-a)$
- Gradient:  $\sigma'(a) = \frac{\exp(-a)}{(1 + \exp(-a))^2} = \sigma(a)(1 - \sigma(a))$

# Classification loss

## Data log-likelihood:

$$\text{loglik}(\mathcal{D}, \mathbf{w}) = \sum_{i=1}^N p(c_i | \phi(\mathbf{x}_i), \mathbf{w}) = \sum_{i=1}^N \underbrace{c_i \log \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)) + (1 - c_i) \log (1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)))}_{\text{loss}_i \dots \text{loss of the } i\text{th sample}}$$

$$\begin{aligned} \frac{\partial \text{loss}_i}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} (c_i \log \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)) + (1 - c_i) \log (1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)))) \\ &=? \end{aligned}$$



# Gradient for Logistic Regression

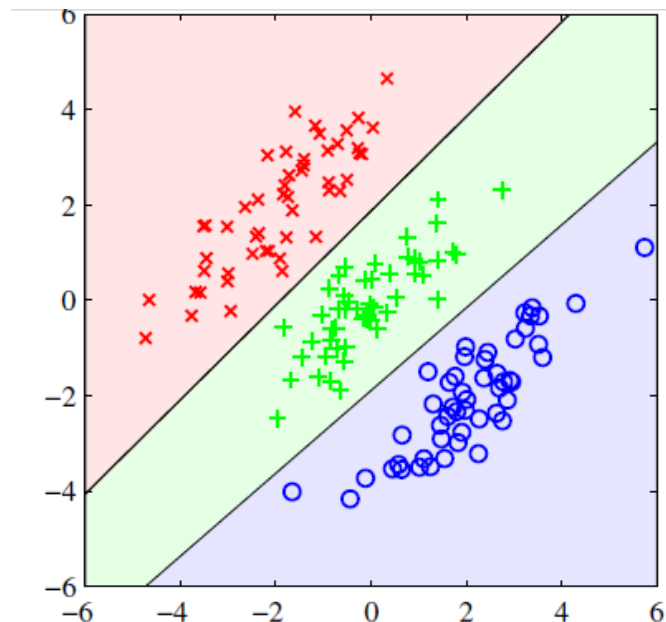
$$\begin{aligned}\frac{\partial \text{loss}_i}{\partial \mathbf{w}} &= \frac{\partial}{\partial \mathbf{w}} \left( c_i \log \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)) + (1 - c_i) \log (1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))) \right) \\ &= c_i \frac{1}{\sigma(\mathbf{w}^T \phi(\mathbf{x}_i))} \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))(1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))) \phi(\mathbf{x}_i) \\ &\quad + (1 - c_i) \frac{1}{1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))} (-) \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))(1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))) \phi(\mathbf{x}_i) \\ &= c_i (1 - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))) \phi(\mathbf{x}_i) - (1 - c_i) \sigma(\mathbf{w}^T \phi(\mathbf{x}_i)) \phi(\mathbf{x}_i) \\ &= (c_i - \sigma(\mathbf{w}^T \phi(\mathbf{x}_i))) \phi(\mathbf{x}_i)\end{aligned}$$

# Multiclass Classification

## Softmax Likelihood function:

$$p(c = i|\mathbf{x}) = \frac{\exp(\mathbf{w}_i^T \phi(\mathbf{x}))}{\sum_{k=1}^K \exp(\mathbf{w}_k^T \phi(\mathbf{x}))}$$

- Each class gets a weight vector
- Higher probability for class  $i$  if  $\mathbf{w}_i^T \phi(\mathbf{x})$  is high
- For  $K = 2$ ,  $\mathbf{w}_2$  is redundant -> better to use sigmoid



# Recap: Multinomial distribution

## Multinomial / Categorical Distribution:

- K different events:  $C \in \{1, \dots, K\}$
- Directly specifies probabilities:  $p(C = k) = \mu_k, \quad \mu_k \geq 0, \quad \sum_{k=1}^K \mu_k = 1$
- Or written with 1-hot-encoding (without an “if” clause)

$$p(c) = \prod_{k=1}^K \mu_k^{h_{c,k}}$$

Depending on the class label of x, selects the correct  $\mu_k$

- where  $\mathbf{h}_x$  is the K-dimensional 1-hot encoding vector, which is one for the dimension  $c = k$  and 0 elsewhere.  $h_{x,k}$  is the k-th element of this vector.
- Think of it as tossing a die

# Multiclass Classification

The multi-class classification problem can be expressed as a **conditional multinomial distribution**:

- I.e. the probability of the event  $c$  depends on the input  $\mathbf{x}$
- We can again use the “exponential trick” to select the correct probability depending on  $c$

$$\begin{aligned} p(c|\mathbf{x}) &= \prod_{k=1}^K p(c = k|\mathbf{x})^{h_{c,k}} \\ &= \prod_{k=1}^K \left( \frac{\exp(\mathbf{w}_k^T \phi(\mathbf{x}_i))}{\sum_{k'=1}^K \exp(\mathbf{w}_{k'}^T \phi(\mathbf{x}_i))} \right)^{h_{c,k}} \end{aligned}$$

# Multiclass Classification

**Data log-likelihood:**  $\text{loglik}(\mathcal{D}, \mathbf{w}_{1:K}) = \sum_{i=1}^N \log p(c_i | \mathbf{x}_i) = \sum_{i=1}^N \underbrace{\sum_{k=1}^K \mathbf{h}_{c_i,k} \log p(k | \mathbf{x}_i)}_{\text{loss}_i \dots \text{loss of the } i\text{th sample}}$

$$= \sum_{i=1}^N \sum_{k=1}^K \mathbf{h}_{c_i,k} \left[ \mathbf{w}_k^T \phi(\mathbf{x}_i) - \log \left( \sum_{j=1}^K \exp(\mathbf{w}_j^T \phi(\mathbf{x}_i)) \right) \right]$$
$$= \sum_{i=1}^N \sum_{k=1}^K \mathbf{h}_{c_i,k} \underbrace{\mathbf{w}_k^T \phi(\mathbf{x}_i) - \log \left( \sum_{j=1}^K \exp(\mathbf{w}_j^T \phi(\mathbf{x}_i)) \right)}_{\text{independent from } k} \underbrace{\sum_k \mathbf{h}_{c_i,k}}_{=1}$$

- Can again be optimized by gradient ascent

# Multiclass Classification

**Gradient:**

$$\frac{\partial \text{loss}_i}{\partial \mathbf{w}_k} = \frac{\partial}{\partial \mathbf{w}_k} \left( \sum_{k=1}^K \mathbf{h}_{c_i, k} \mathbf{w}_k^T \phi(\mathbf{x}_i) - \log \left( \sum_{j=1}^K \exp(\mathbf{w}_j^T \phi(\mathbf{x}_i)) \right) \right)$$

=?

# Takeaway messages

## What have we learned today?

- Refresher on probability theory and maximum likelihood
- Relation between maximum likelihood and least squares
- What is a linear classification problem ...
- ... and how to formalize it as likelihood maximization problem
  - Sigmoid likelihood for binary classification
  - Soft-max likelihood for multi-class
- What is gradient descent, stochastic gradient descent and mini-batches?
- How to apply gradient descent to logistic regression











