

Model Selection

Maschinelles Lernen 1 -
Grundverfahren WS20/21

Prof. Gerhard Neumann
KIT, Institut für Anthropomatik und Robotik

Learning Outcomes

- Understand the overfitting problem and ...
- ... its relation to the complexity of the model class
- Bias variance tradeoff
- Why we need test-sets and cross-validation
- Understand different regularization methods

Agenda for today

Model Selection



- Overfitting and model complexity
- Bias variance trade-off
- Hold-out set and cross validation

Regularization:

- Limit complexity
- Penalty terms
- Early stopping
- Noise and data augmentation

Zoo of algorithms for machine learning

- **Regression: Continuous output labels**
 - Linear regression, Polynomial Regression, kNN, Regression Trees, Gaussian Processes, Neural Nets
- **Classification: Discrete / Nominal output labels**
 - Logistic Regression, Decision Trees, Neural Nets, SVMs, kNN



Model Complexity

For most of these algorithms, we have to choose the model complexity

- **Linear Regression:** number of features, regularization coefficient
- **Decision Trees:** maximum depth, number of leaves
- **Neural Networks:** number of layers, number of neurons
- **Support Vector Machine:** which features, regularization
- **Gaussian Processes:** kernel bandwidth

Choosing the right complexity is a model selection problem!

- And one of the most fundamental problems in ML

True risk vs. empirical risk

True risk: performance on a random test point (x,y)

- **Classification:** probability of misclassification
- **Regression:** expected squared error

$$p(y \neq f(x))$$

$$\mathbb{E}_{\mathbf{x},y}[(f(\mathbf{x}) - y)^2]$$

➤ **Unknown!**

Empirical risk: performance on the training set

- **Classification:** proportion of misclassified samples
- **Regression:** average squared error

$$\frac{1}{n} \sum_i \mathbb{I}(f(\mathbf{x}_i) \neq y_i)$$
$$\frac{1}{n} \sum_i (f(\mathbf{x}_i) - y_i)^2$$

➤ **Can be evaluated**

Overfitting


Is the following predictor good?

$$f(x) = \begin{cases} y_i, & \text{if } x = x_i \text{ for } i = 1 \dots n \\ \text{any other value,} & \text{else} \end{cases}$$

Empirical risk? Zero!

True risk? Huge!

- Will predict poorly on unseen data
- Large generalization error





Model Selection

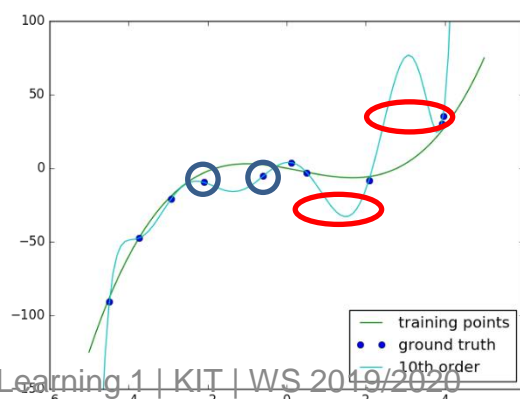
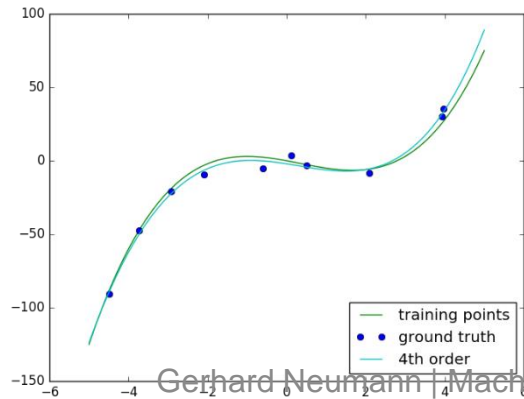

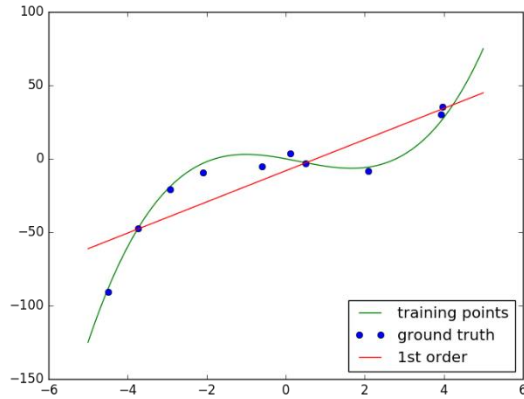
Choose complexity of the model class

- **Too complex: Overfitting**
 - Fit noise in the data
 - Unspecified behavior between data points
 - Not enough data
- **Too simple: Underfitting**
 - We can not represent the underlying function

Lets look at an example...



Polynomial regression up to kth order



Last model overfits the data:

- Fits the noise
- Unspecified behavior in between datapoints


Effect of Model Complexity

Overfitting:

Small empirical risk, but true risk is high


Underfitting:

High empirical risk and true risk



Bias-Variance Decomposition

Regression example: $y = f(x) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$



$$\text{Expected Loss} = \text{Variance} + \text{Bias}^2 + \text{Noise}$$

Bias:


- Due to restriction of your model class
- Also called “structure error”

Variance:


- Due to randomness of the data set

Noise:

- Nothing we can do about it...



Understanding the true risk



Regression example: $y = f(x) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$

Expected Loss:


$$R(\hat{f}_{D_n}) = \mathbb{E}_{D_n} \left[\underbrace{\mathbb{E}_{x,y} \left[(\hat{f}_{D_n}(x) - y)^2 \right]}_{\text{Expected error for training with a specific dataset}} \right]$$

Expected error for training with a specific dataset
Expectation is done w.r.t to all possible inputs and corresponding outputs

- \hat{f}_{D_n} is the estimate of f we obtain using data D_n
- The expectation is done w.r.t all training-sets D_n of size n . What does that mean?
 - Assume there is a process that can generate data sets n data-points, e.g.:
 - Sample x_i uniformly in range $[-1, 1]$
 - Sample y_i as $y_i = f(x_i) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$
 - Repeat n times
 - The expectation averages the error over all training-sets D_n

Understanding the true risk

Regression example: $y = f(\mathbf{x}) + \epsilon$, $\epsilon \sim \mathcal{N}(0, \sigma^2)$



Expected Loss:

$$\begin{aligned} R(\hat{f}_{D_n}) &= \mathbb{E}_{D_n} \left[\mathbb{E}_{x,y} \left[(\hat{f}_{D_n}(\mathbf{x}) - y)^2 \right] \right] \\ &= \underbrace{\mathbb{E}_{D_n} \left[\mathbb{E}_{x,y} \left[(\hat{f}_{D_n}(\mathbf{x}) - \hat{f}_*(\mathbf{x}))^2 \right] \right]}_{\text{Variance}} + \underbrace{\mathbb{E}_{x,y} \left[(\hat{f}_*(\mathbf{x}) - f(\mathbf{x}))^2 \right]}_{\text{Bias}^2} + \underbrace{\sigma^2}_{\text{noise}} \end{aligned}$$

- $\hat{f}_*(\mathbf{x}) = \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})]$ is the average estimate, averaged over all data sets of size n
- ... which can approximately be seen as training with infinite data

Bias-Variance Decomposition

Observations:


- We can not get better than the noise

Bias: $\mathbb{E}_{x,y} \left[\left(\hat{f}_*(x) - f(x) \right)^2 \right]$

- Difference of true function to the “best” estimate
- The best you can do with your model class
- Also called “approximation error”

Variance: $\mathbb{E}_{D_n} \left[\mathbb{E}_{x,y} \left[\left(\hat{f}_{D_n}(x) - \hat{f}_*(x) \right)^2 \right] \right]$

- Difference of the estimates to the “best” estimate
- Due to limited size of the data set
- Depends on number of data-points
- Also called “estimation error”




Derivation of the Bias-Variance Trade-off

A lot of math... look in the appendix if interested

Bias-Variance Trade-off

Each point represents a fitted model with a different dataset

- **Low variance, high bias:** Underfitting
- **High variance, low bias:** Overfitting
- **High variance, high bias:** something is terribly wrong
- **Low variance, low bias:** too good to be true



Agenda for today

Model Selection

- Overfitting and model complexity
- Bias variance trade-off
- **Hold-out set and cross validation**

Regularization:

- Limit complexity
- Penalty terms
- Early stopping
- Noise and data augmentation

Evaluation Methods

We have seen that the empirical risk on the training set is not a good indicator for the quality of your model. What can we do?

- Hold-out method
- Cross-validation

Hold-out method

- We would like to pick the model M with the smallest generalization error
- Can judge the generalization error by independent data set (not used for training)

Hold-out procedure: n datapoints available $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$

1. Split into 2 datasets:

Training Data

$$D_T = \{(\mathbf{x}_i, y_i)\}_{i=1}^m$$

Validation Data

$$D_V = \{(\mathbf{x}_i, y_i)\}_{i=m+1}^n$$

2. Train on training data to obtain $\hat{f}_{D_T}(\mathbf{x})$ for each model class M

3. Evaluate resulting estimators on validation data, e.g: $\text{MSE}(D_V, \hat{f}_{D_T}) = \frac{1}{n-m} \sum_{i=m+1}^n (\hat{f}_{D_T}(\mathbf{x}_i) - y_i)^2$

4. Pick model with best validation loss

Hold-out method

Drawbacks:

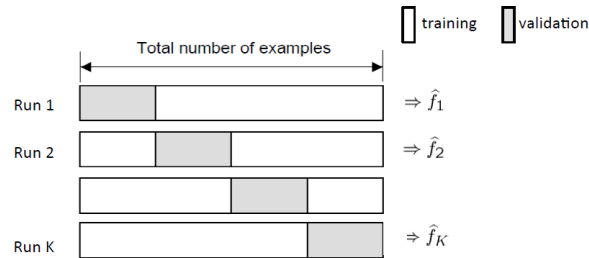
- Costly in terms of data
- “Unlucky” splits might give misleading results

Cross-validation methods fix these issues

Cross validation

K-fold cross validation

1. Create k-fold partition of the dataset
2. Estimate k hold-out predictors using 1 partition as validation and k-1 partitions as training set




k predictors for each model class:

Cross validation

Leave-One-Out (LOO) cross validation

1. Special case with $k = n$
2. Consequently, estimate n hold-out predictors using 1 sample as validation and $n-1$ samples as training set

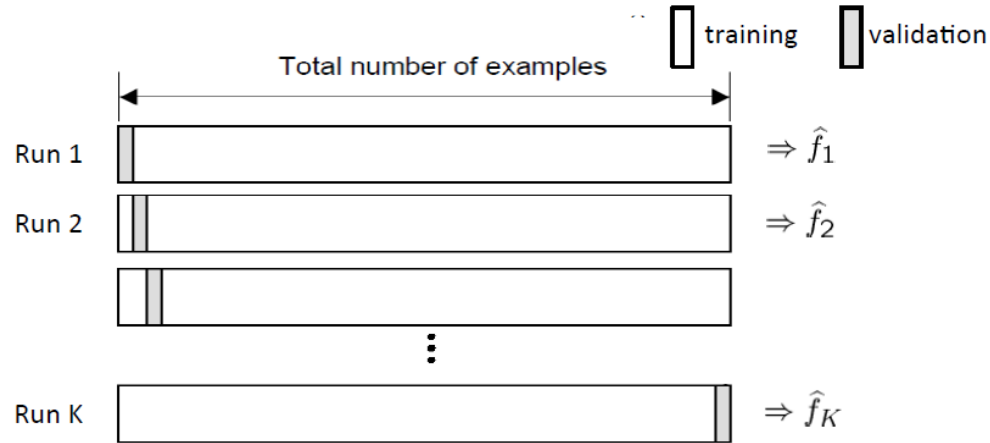


k predictors for each model class:

Cross validation

Random sub-sampling

1. Randomly sample a fraction of $\alpha * n$ ($0 < \alpha < 1$) data points for validation
2. Train on remaining points and validate, repeat K times



k predictors for each model class:

Regularization techniques

How to avoid overfitting?

- Limit the complexity of the model (# neurons, # of leaves, etc...)
- Regularization penalty
- Early stopping
- Noise and Data-Augmentation
- ...

Occam's Razor

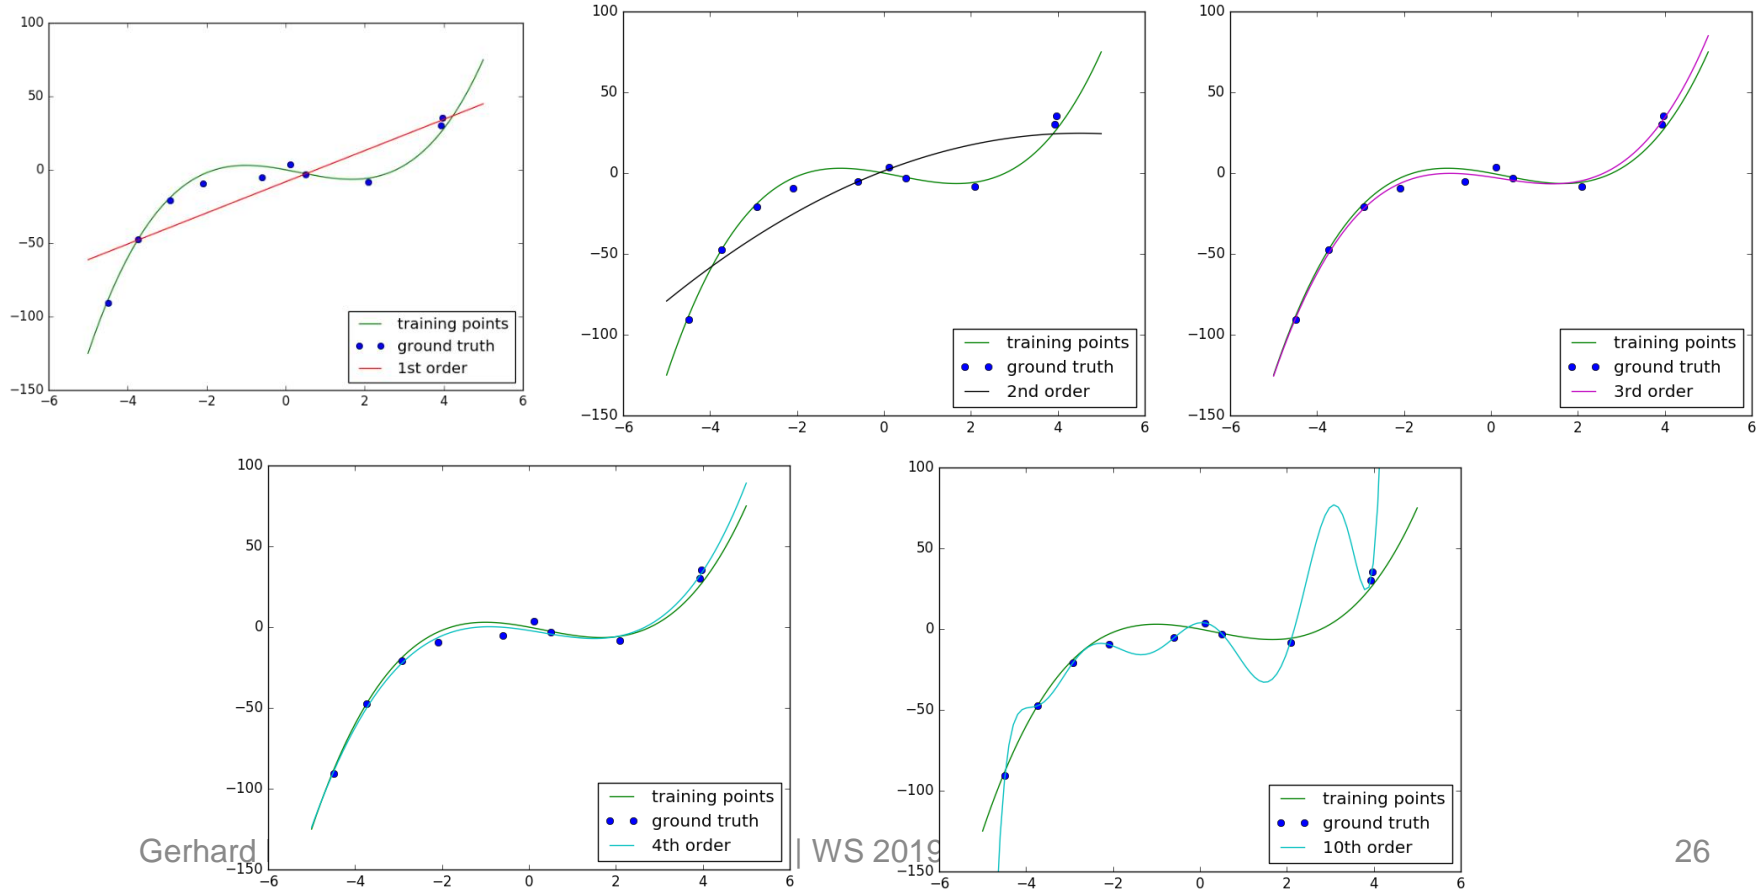
- Named after William of Occam – AD 1300s
- Prefer **simpler explanations** over more complex ones
 - “Numquam ponenda est pluralitas sine necessitate”
 - (Never posit plurality without necessity.)
- Historically, a widely prevalent idea across different schools of philosophy
- Directly applicable for model selection in ML

Occam's Razor




“When faced with two equally good hypotheses, always choose the simpler.”

Limit complexity – Example with polynomials



Model Selection for polynomial regression

- **Overfitting:**
 - Training error goes down
 - Validation error goes up
- **Underfitting:**
 - Training + Validation error are high
- **Optimum:** 3rd or 4th degree



Regularization penalty

Can be used for most optimization-based algorithms

- Linear Regression + Classification, Neural Networks, GPs, ...

We, typically optimize a (sample-based) **loss plus a regularization penalty**

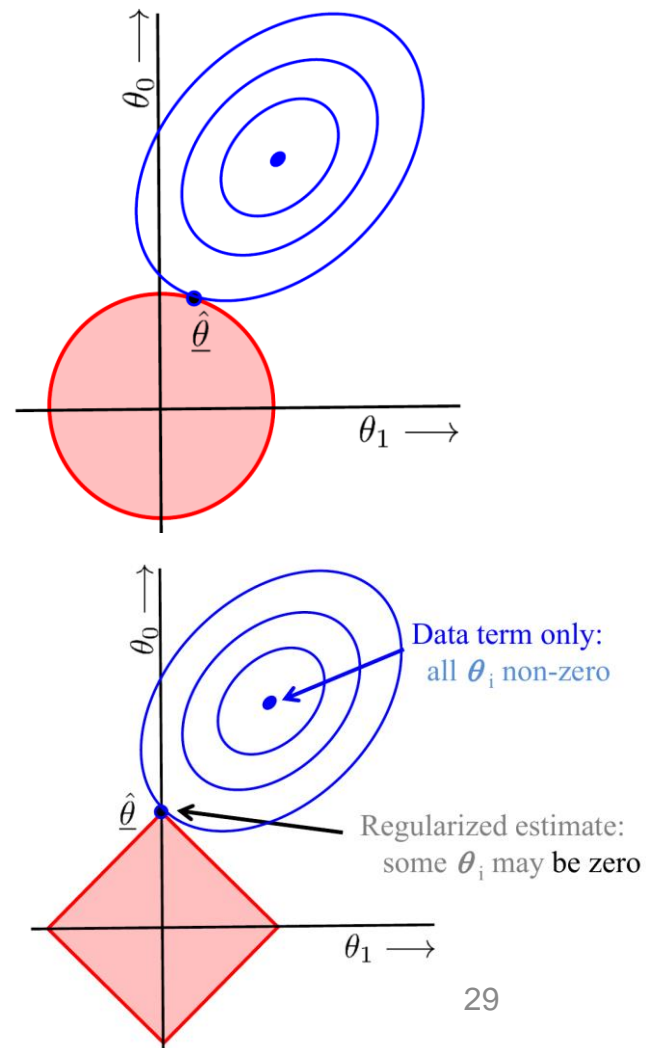
$$\arg \min_{\text{parameters } \boldsymbol{\theta}} \sum_{i=1}^N l(\mathbf{x}_i, \boldsymbol{\theta}) + \lambda \text{penalty}(\boldsymbol{\theta})$$

- Penalty keeps parameters small
- Small parameters -> **smoother function** estimate
- Implicitly **limits the complexity** of the learned model (larger lambda -> smaller complexity)

Regularization penalty


Which penalty functions can we use?

- **L_2 penalty:** $\text{penalty}(\boldsymbol{\theta}) = \|\boldsymbol{\theta}\|_2 = \sum_d \theta_d^2$
 - Easy to optimize (strongly convex)
 - Closed form solutions exists
 - Redundant parameters will be close to 0, but never 0
- **L_1 penalty:** $\text{penalty}(\boldsymbol{\theta}) = \|\boldsymbol{\theta}\|_1 = \sum_d |\theta_d|$
 - Induces **sparse** solutions
 - Called “Lasso” regularization
 - Much harder to optimize (not in this lecture)



Example: ridge regression

Ridge regression with polynomial of degree $n=15$

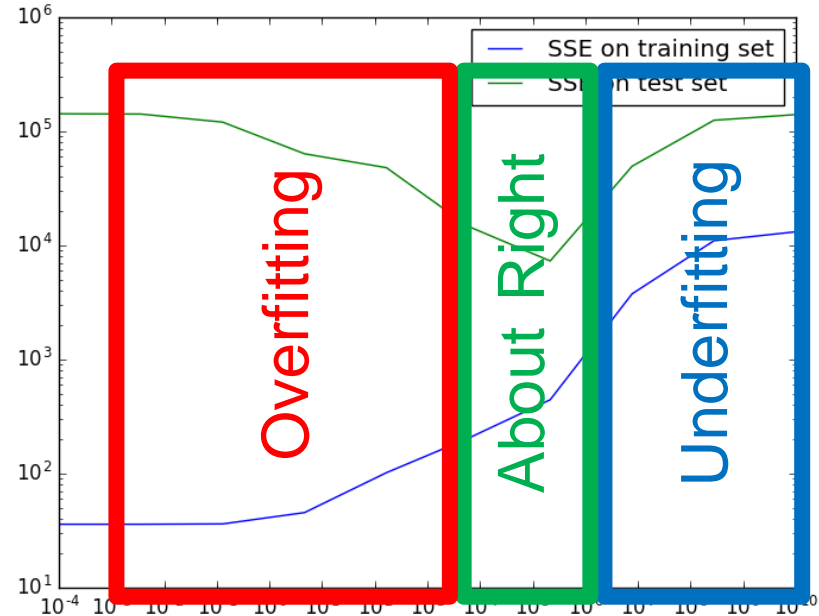


Influence of the regularization constant

Example: Ridge regression

Influence of the lambda parameter

- High Lambda: Underfitting
 - High training and validation error
- About Right:
 - Validation error is minimal
- Small Lambda: Overfitting
 - Small training error, but high validation error



Early stopping

Idea: don't train to too small training error

- Used with incremental learning rules (e.g. gradient descent)
- Prevent overfitting: do not push the model too much; use validation error to decide when to stop
- Implicitly limits complexity

Early stopping

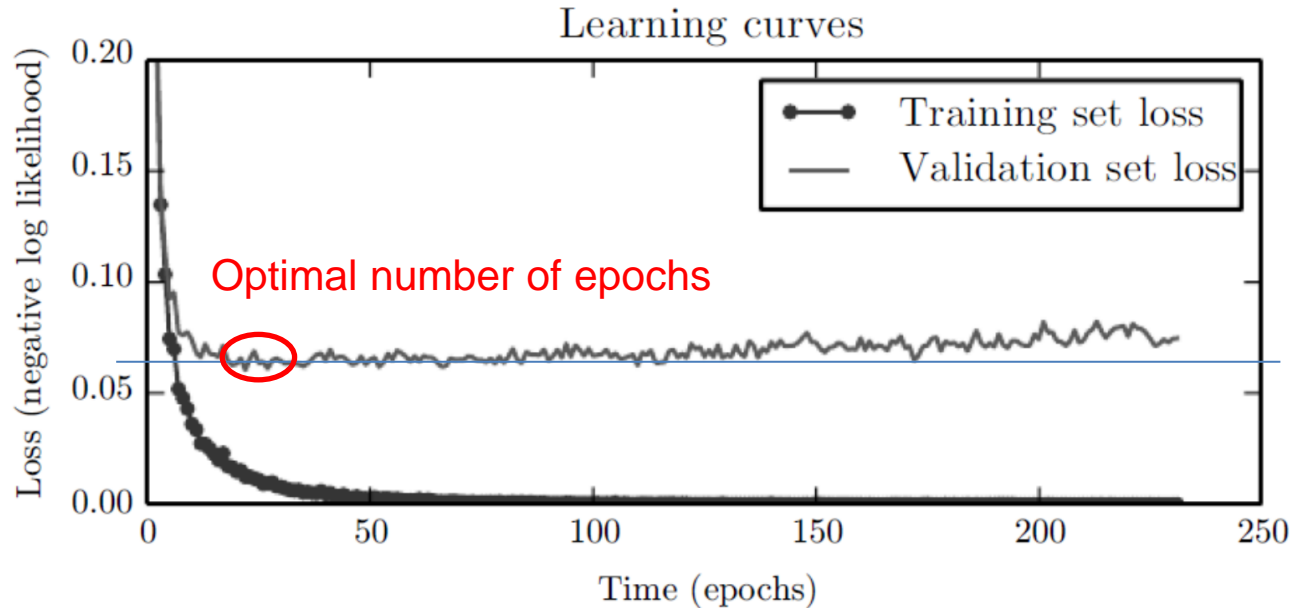


Figure from *Deep Learning*,
Goodfellow, Bengio and Courville

- Validation **error goes up due to overfitting**

Early Stopping

- When training, also output validation error
- Every time validation error improved, store a copy of the weights
- When validation error not improved for some time, stop
- Return the copy of the weights stored

Early stopping as regularizer

- Early stopping has similar effects than L_2 regularization

Advantage

- Efficient: along with training; only store an extra copy of weights
- Simple: no change to the model/algo
- No hyper-parameter (such as λ)

Disadvantage

- need validation data





Figure from *Deep Learning*,
Goodfellow, Bengio and Courville


Robustness to noise

All solutions are valid according to the classification loss



Adding noise to the inputs

- Rules out unlikely / not robust solutions



Equivalence to L2 regularization

For a linear regression model, the noise model is given by

$$f(\mathbf{x} + \boldsymbol{\epsilon}) = \mathbf{w}^T (\mathbf{x} + \boldsymbol{\epsilon}), \quad \boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \lambda \mathbf{I})$$

This leads to the following loss:

$$\begin{aligned} \text{MSE}(\mathbf{w}) &= \mathbb{E}_{\mathbf{x}, y, \boldsymbol{\epsilon}} [(\mathbf{w}^T \mathbf{x} - y + \mathbf{w}^T \boldsymbol{\epsilon})^2] \\ &= \underbrace{\mathbb{E}_{\mathbf{x}, y} [(\mathbf{w}^T \mathbf{x} - y)^2]}_{n\text{SSE}(\mathbf{w})} + 2 \underbrace{\mathbb{E}_{\mathbf{x}, y, \boldsymbol{\epsilon}} [(\mathbf{w}^T \mathbf{x} - y) \mathbf{w}^T \boldsymbol{\epsilon}]}_{=0, \text{zero mean, i.i.d. noise}} + \underbrace{\mathbb{E}_{\mathbf{x}, y, \mathbf{w}} [(\mathbf{w}^T \boldsymbol{\epsilon})^2]}_{\lambda \mathbf{w}^T \mathbf{w}} \\ &= n\text{SSE}(\mathbf{w}) + \lambda \|\mathbf{w}\|_2^2 \end{aligned}$$

I.e., for linear regression, **input noise** is the same as **L2 regularization**

- For other models, the effect is similar, but not exactly the same

Data Augmentation

Create additional artificial samples

Yet, be careful about the transformation applied

- Example: classify 'b' and 'd'
- Example: classify '6' and '9'

Horizontal Flip



Crop



Rotate



Takeaway messages

What have we learned today?

- **Never** use training set to evaluate your model!
- Understand the causes of overfitting
 - Learn noise in the data
 - Unspecified behaviour between data points
- Bias-Variance tradeoff and its relation to the model complexity
- How to evaluate models
- Different regularization strategies



Self-test questions

What you should understand by now:

- Why is it a bad idea to evaluate your algorithm on the training set
- What is the difference between true and empirical risk
- The true risk can be decomposed in which parts?
- How is the bias and the variance of a learning algorithm defined and how do they contribute to the true risk?
- What is the advantage/disadvantage of k-fold CV vs. the Hold-out method?
- Why does it make sense to penalize the norm of the weight vector?
- Which norms can we use and what are the different effects?
- What is the effect of early stopping?

Appendix: Derivation of the Bias-Variance Trade-off 1/2

$$\begin{aligned} R(\hat{f}_{D_n}) &= \mathbb{E}_{x,y,D_n} \left[(\hat{f}_{D_n}(\mathbf{x}) - y)^2 \right] = \mathbb{E}_{x,y,D_n} \left[((\hat{f}_{D_n}(\mathbf{x}) - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})]) + (\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y))^2 \right] \\ &= \mathbb{E}_{x,y,D_n} \left[(\hat{f}_{D_n}(\mathbf{x}) - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})])^2 + (\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y)^2 \right. \\ &\quad \left. + 2(\hat{f}_{D_n}(\mathbf{x}) - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})])(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y) \right] \\ &= \mathbb{E}_{x,y,D_n} \left[(\hat{f}_{D_n}(\mathbf{x}) - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})])^2 \right] + \mathbb{E}_{x,y,D_n} \left[(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y)^2 \right] \\ &\quad + 2\mathbb{E}_{x,y} \left[\underbrace{(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})])}_{=0} (\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y) \right] \end{aligned}$$

Appendix: Derivation of the Bias-Variance Trade-off 2/2

$$R(\hat{f}_{D_n}) = \underbrace{\mathbb{E}_{x,y,D_n} \left[(\hat{f}_{D_n}(\mathbf{x}) - \mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})])^2 \right]}_{\text{variance}} + \mathbb{E}_{x,y,D_n} \left[(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y)^2 \right]$$

2nd term:

$$\begin{aligned} \mathbb{E}_{x,y,D_n} \left[(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - y)^2 \right] &= \mathbb{E}_{x,y} \left[(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - f(\mathbf{x}) - \epsilon)^2 \right] \\ &= \underbrace{\mathbb{E}_{x,y} \left[(\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - f(\mathbf{x}))^2 \right]}_{\text{bias}^2} + \underbrace{\mathbb{E}_{x,y} [\epsilon^2]}_{\text{noise}} - 2 \underbrace{\mathbb{E}_{x,y} \left[\epsilon (\mathbb{E}_{D_n}[\hat{f}_{D_n}(\mathbf{x})] - f(\mathbf{x})) \right]}_{0 \text{ due to zero mean i.i.d noise}} \end{aligned}$$

