

Nearest Neighbour Algorithms, Trees and Forests

Maschinelles Lernen 1 -
Grundverfahren WS20/21

Prof. Gerhard Neumann
KIT, Institut für Anthropomatik und Robotik

Learning Outcomes

What will we learn today?

Nearest Neighbors:

- What an **non-parametric/instance** based learning algorithm is
- ... start with the most simple non-parametric algorithm: k-Nearest Neighbor
- What is the **curse of dimensionality**
- How to **compute** the nearest neighbors **efficiently**

Trees:

- How can we use trees for classification and regression
- Why should we use ensembles of trees (forests)?
- Why should these be random to some extend?

Today's Agenda!

Nearest Neighbour Algorithms:

- k-Nearest Neighbour Classifiers
- Curse of dimensionality
- Indexing with KD-trees

Tree-based methods

- **For regression:** Regression Tree
- **For classification:** Decision Tree
- Almost the same algorithms!

Random forest

- Bagging predictors
- Randomization

K-Nearest Neighbor Algorithms

Non-parametric Methods

Non-parametric methods **store all the training data** and use the training data for doing predictions. They **do not adapt parameters** or a parametric model. They are also often referred to as **instance-based methods**.

- ✓ Complexity adapts to training data
- ✓ Very fast at training
- × Slow for prediction
- × Hard to use for high-dimensional input

Algorithms:

- k-Nearest Neighbor Algorithm (today)
- Locally Weighted Regression (not covered)
- Kernel Methods and Gaussian Processes (later)

K-Nearest Neighbour Classifier

To classify a new input vector x , examine the **k-closest training data points** to x and assign the object to the most frequently occurring class


- common values for k : $k = 3$, $k = 5$

When to consider:

- Can measure distances between data-points
- Less than 20 attributes per instance
- Lots of training data

Advantages:


- Training is very fast
- Learn complex target functions
- Similar algorithm can be used for regression



Decision Boundaries

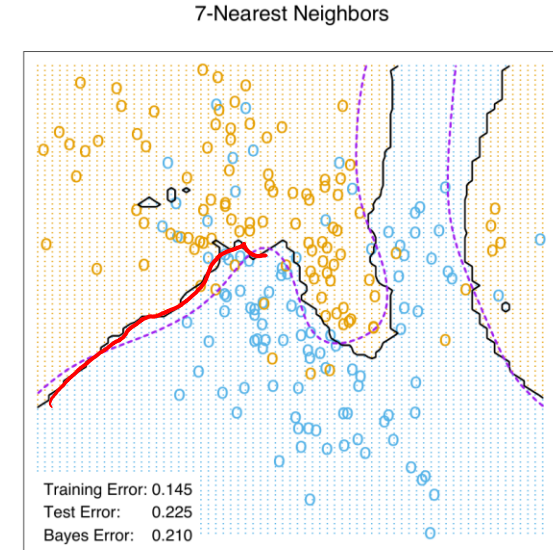
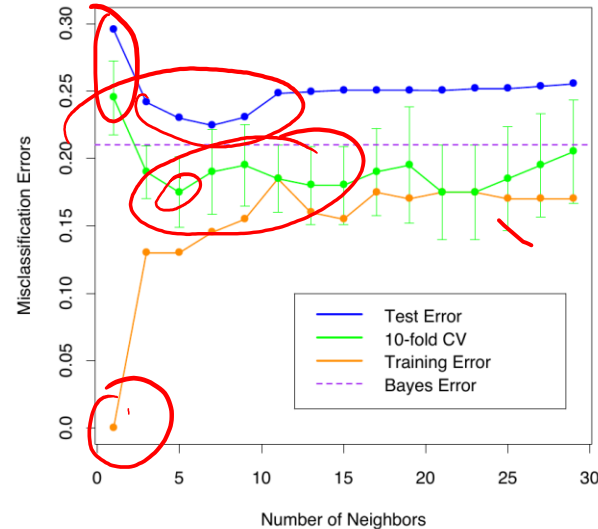
- The nearest neighbour algorithm does not explicitly compute decision boundaries.
- However, the decision boundaries form a subset of the Voronoi diagram for the training data.
- The **more data points** we have, the **more complex** the **decision boundary** can become

1-NN Decision Surface



Example Result

- Bayes error: error of perfect decision boundary
- Increasing k reduces variance, increases bias
 - $K < 7$: overfitting
 - $K > 7$: underfitting
- Has to be selected by cross-validation



Distance Metrics

Most common distance metric is **Euclidean distance (ED)**:

$$d(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\| = \sqrt{\left(\sum_{k=1}^d (\mathbf{x}_k - \mathbf{y}_k)^2 \right)}$$

- ED makes sense when different features are commensurate; each is variable measured in the same units.
- If the units are different, say length and weight, **data needs to be normalized**:

$$\tilde{\mathbf{x}} = (\mathbf{x} - \boldsymbol{\mu}) \oslash \boldsymbol{\sigma}$$

- Mean $\boldsymbol{\mu}$, standard deviation $\boldsymbol{\sigma}$, element-wise division \oslash
- I.e. resulting input dimensions are zero mean, unit variance

Distance Metrics

- **Cosine Distance:** Good for documents, images, etc.

$$d(\mathbf{x}, \mathbf{y}) = 1 - \frac{\mathbf{x}^T \mathbf{y}}{\|\mathbf{x}\| \|\mathbf{y}\|}$$

- **Hamming Distance:** For string data / categorical features

$$d(\mathbf{x}, \mathbf{y}) = \sum_{k=1}^d (\mathbf{x}_k \neq \mathbf{y}_k)$$

- **Manhattan Distance:** Coordinate-wise distance

$$d(\mathbf{x}, \mathbf{y}) = \sum_{k=1}^d |\mathbf{x}_k - \mathbf{y}_k|$$

Distance Metrics

- **Mahalanobis Distance:** Normalized by the sample covariance matrix – unaffected by coordinate transformations.

$$d(\mathbf{x}, \mathbf{y}) = \|\mathbf{x} - \mathbf{y}\|_{\Sigma^{-1}} = \sqrt{(\mathbf{x} - \mathbf{y})^T \Sigma^{-1} (\mathbf{x} - \mathbf{y})}$$


k-NN and irrelevant features

- No irrelevant input:



- Class can be clearly determined

- Added irrelevant dimension:




- Neighborhood needs to be increased
- Heavily affected by noise
- Needs much more training data

The performance of k-NN degrades with more (irrelevant) dimensions


Curse of dimensionality

Example 1: What fraction of the points in a cube lie outside the sphere inscribed in it?



$$\frac{4r^2 - \pi r^2}{4r^2}$$

$$1 - \frac{\pi r^2}{4r^2} \\ = 1 - \frac{\pi}{4}$$



$$1 - \frac{4/3\pi r^3}{8r^3} \\ = 1 - \frac{\pi}{6}$$

- For $d \rightarrow \infty$ this fraction approaches 1!

Curse of dimensionality

Most of the points in high dimensional spaces are far away from the origin!

- In 2 or 3 dimensions, most points are near the center
- Need more data to “fill up the space”

Bad news for nearest neighbor classification in high dimensional spaces

- Even if most/all features are relevant, in high dimensional spaces, most points are equally far from each other!
- “Neighborhood” becomes very large

Remedies (to some extend):

- Most “real-world” data is not uniformly distributed in the high dimensional space
- E.g.: Dimensionality reduction techniques, manifold learning
- Feature selection (pick a good set based on a validation set)


Finding the neighbours: KD-Trees

Problem: given a sample set $S = \{x_1, \dots, x_N\}$, find the k-NNs of test point x^* .

Building the tree: for each non-leaf node

- Choose dimension (e.g., longest hyperrectangle).
- Choose median as pivot
- Split node according to (pivot, dimension).

Balanced tree, binary space partitioning.




KD-Trees

Finding the neighbours ($k = 1$):

- Find region containing \mathbf{x} (starting from root node, move to child node based on node test).
- Save region point $\mathbf{x}^* = \mathbf{x}_0$ as current best.
- Move up tree and recursively search regions intersecting hypersphere $S(\mathbf{x}, \|\mathbf{x} - \mathbf{x}^*\|)$
- Update \mathbf{x}^* if new nearest neighbour has been found

For $k > 1$:

- Same algorithm, but save \mathbf{x}^* as k-nearest neighbour.
- **Complexity:** $O(k \log N)$



k-NN Summary

Probably the oldest and simplest learning algorithm

- Prediction is expensive.
- Efficient data structures help. k-D trees: the most popular, works well in low dimensions
- Good baseline: If you do not beat k-NN, you are doing something wrong

Requires a distance measure between instances

- Partitions the space into a Voronoi Diagram
- Beware the curse of dimensionality



Trees and Forests

Regression and Classification Trees

Grow a binary tree


- At each node, “split” the data into two “daughter” nodes.
- Splits are chosen using a splitting criterion.
- Bottom nodes are “terminal” nodes.

For regression:

- the predicted value at a node is the **average response** variable for all observations in the node.

For classification:


- the predicted class is the **most common class** in the node (majority vote).
- Can also get estimated probability of membership in each of the classes



A classification tree

Predict hepatitis (0=absent, 1=present)
using protein and alkaline phosphate.

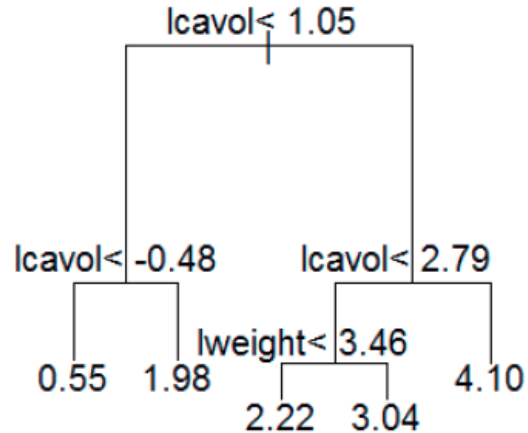

“Yes” goes left.



A regression tree

Predict (log) prostate specific antigen from

- Log cancer volume
- Log prostate weight



Splitting criterion

- **Regression:** Minimum residual sum of squares

$$\text{RSS} = \sum_{\text{left}} (y_i - \bar{y}_L)^2 + \sum_{\text{right}} (y_i - \bar{y}_R)^2$$

- where \bar{y}_L and \bar{y}_R are the average label values in the left and right subtree
- Split such that variance in subtrees is minimized

Splitting criterion (thats the second mathy slide...)

- **Classification:** Minimum entropy in subtrees

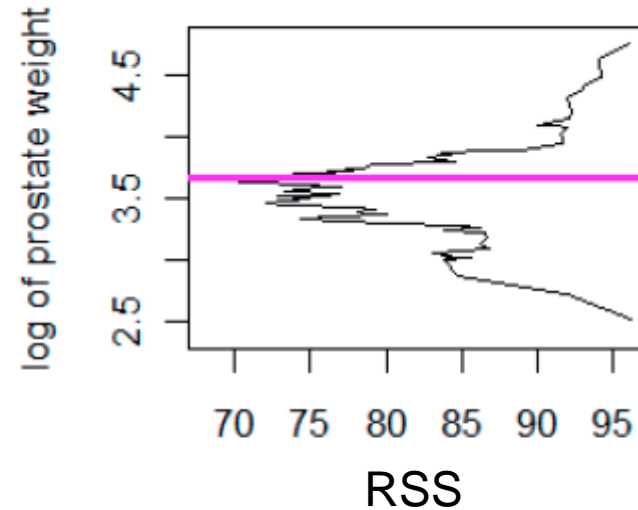

$$\text{score} = N_L H(p_L) + N_R H(p_R)$$

- where $H(p_L) = - \sum_k p_L(k) \log p_L(k)$ is the entropy in the left sub-tree
- and $p_L(k)$ is the proportion of class k in left tree

- **Entropy is a measure of uncertainty**



- Split such that class-labels in sub-trees are „pure“

Finding the best horizontal split





Best horizontal split is at 3.67 with $RSS = 68.09$.

Finding the best vertical split




Best vertical split is at 1.05 with $RSS = 61.76$.

Creating the root node





Finding the best split in the left node





Best horizontal split is at 3.66 with $RSS = 16.11$.

Finding the best split in the left node





Best vertical split is at -0.48 with $RSS = 13.61$.

Building the regression tree...


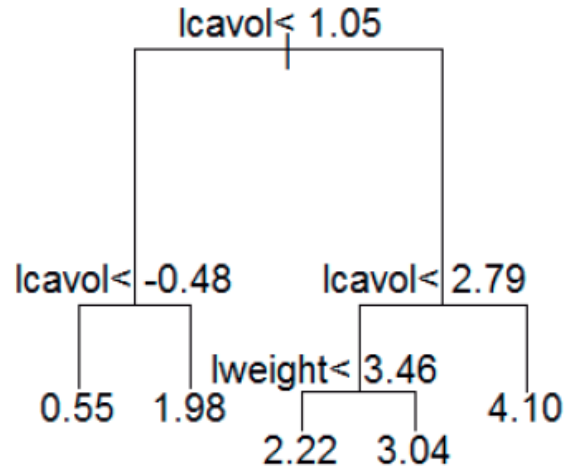
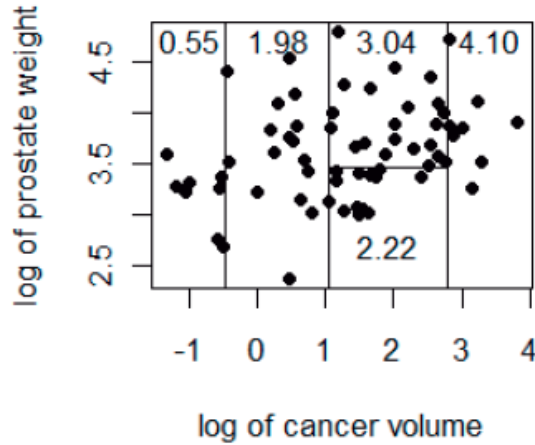


Finding the best split in the right node...




Best horizontal split is at 3.07 with $RSS = 27.15$.


Skipping some steps... final result




Classification Tree




Classification Tree






Classification Tree



Classification Tree



When do we stop?

There are many stopping criterias, the 2 main ones are:

Stop if:


- Minimum number of samples per node
 - Maximum depth
- ... has been reached

Both criterias again influence the **complexity** of the tree !

Controlling the tree complexity

Small number of samples per leaf:


- Tree is very sensitive to noise



Controlling the tree complexity

Small number of samples per leaf:

- Tree is very sensitive to noise




Controlling the tree complexity

Small number of samples per leaf:

- Tree is very sensitive to noise

Large number of samples per leaf:


- Tree not expressive enough



Model-Selection for Regression Trees

Evaluate error on validation-set

- Overfitting for $\text{min_samples} = 1$
- Underfitting for $\text{min_samples} > 2$
- Larger min_samples \rightarrow lower complexity



Classification and Regression Trees

Advantages

- Applicable to both regression and classification problems.
- Handle categorical predictors naturally.
- Computationally simple and quick to fit, even for large problems.
- No formal distributional assumptions
- Can handle highly non-linear interactions and classification boundaries.
- Automatic variable selection.
- Very easy to interpret if the tree is small.

Classification and Regression Trees (CART)


Disadvantages

- *Accuracy* - current methods, such as NNs, support vector machines and ensemble classifiers often have 30% lower error rates than CART.
- *Instability* – if we change the data a little, the tree picture can change a lot. So the interpretation is not as straightforward as it appears.


Nowadays, we can do better! **Random Forests!**

Key Idea: Use multiple trees to improve accuracy


Data and Underlying Function




Single Regression Tree



10 Regression Trees




Average of 100 Regression Trees




Key Idea: Use multiple trees to improve accuracy


Hard problem for a single tree:



Single tree:



25 Averaged Trees:

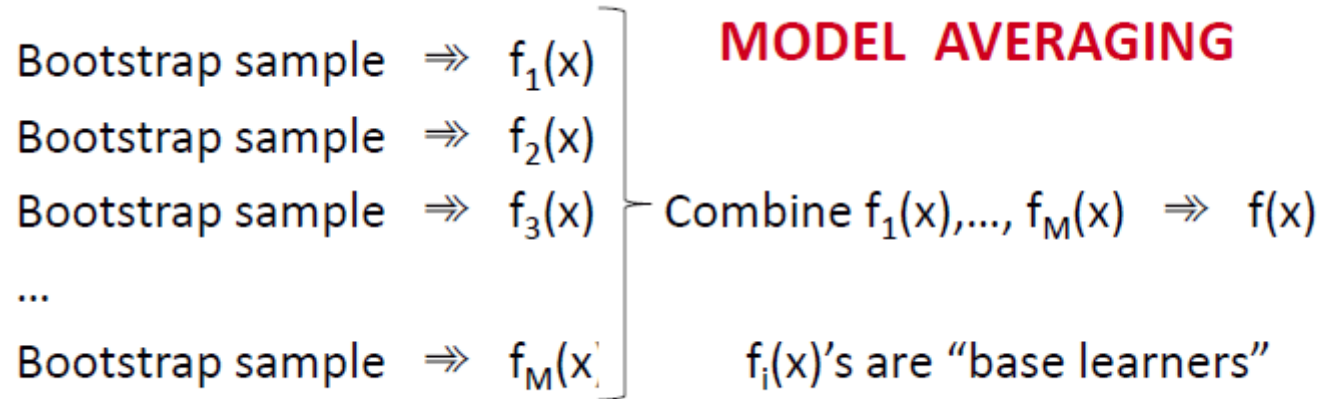


How do we get **variability** in the **trees**?

Bagging (Bootstrap Aggregating)

Breiman, “Bagging Predictors”, *Machine Learning*, 1996.

Fit classification or regression models to **bootstrap samples** from the data and combine by **voting** (classification) or **averaging** (regression).



Bagging (Bootstrap Aggregating)

A **bootstrap sample** is chosen at **random *with replacement*** from the data. Some observations end up in the bootstrap sample more than once, while others are not included (“out of bag”).

Variance reduction

In general:
$$\text{Var} \left[\frac{1}{M} \sum_{i=1}^M X_i \right] = \frac{1}{M^2} \text{Var} \left[\sum_{i=1}^M X_i \right] = \frac{1}{M} \text{Var} [X], \quad \text{if } X \text{ i.i.d.}$$

- i.e., ideally, the variance would reduce linearly with the number of trees

In practice:
$$\text{Var} \left[\frac{1}{M} \sum_{i=1}^M \text{Tree}_i \right] > \frac{1}{M} \text{Var} [\text{Tree}], \quad \text{as trees are still correlated}$$

- But variance reduction is still significant
- Bagging reduces the **variance** of the base learner but has limited effect on the **bias**
 - I.e. no overfitting: The more trees the better
- It's most effective if we use *strong* base learners that have very little bias but high variance (unstable). E.g. trees.

Bagging CART

| Dataset | # cases | # vars | # classes | CART | Bagged CART | Decrease % |
|---------------|---------|--------|-----------|------|-------------|------------|
| Waveform | 300 | 21 | 3 | 29.1 | 19.3 | 34 |
| Heart | 1395 | 16 | 2 | 4.9 | 2.8 | 43 |
| Breast Cancer | 699 | 9 | 2 | 5.9 | 3.7 | 37 |
| Ionosphere | 351 | 34 | 2 | 11.2 | 7.9 | 29 |
| Diabetes | 768 | 8 | 2 | 25.3 | 23.9 | 6 |
| Glass | 214 | 9 | 6 | 30.4 | 23.6 | 22 |
| Soybean | 683 | 35 | 19 | 8.6 | 6.8 | 21 |

Leo Breiman (1996) “Bagging Predictors”, Machine Learning, 24, 123-140.

Randomization

Grow a **forest** of many trees. (R default is 500)

Grow each tree on an independent **bootstrap sample** from the training data.

- Sample N cases at random with replacement.

At each node:

1. Select m variables **at random** out of all M possible variables (independently for each node).
2. Find the best split on the selected m variables.

Grow the trees to maximum depth (classification).

Vote/average the trees to get predictions for new data.


Why does that work?

Intuition: Why randomization?


- Increase variability of the single trees
- A single tree is less likely to over-specialize
- The trees are less likely to overfit

Random Regression Forests


Num Trees = 1



Num Trees = 10



Num Trees = 50
num_trees = 50



We can represent almost continuous functions!

Random Forests

| Dataset | # cases | # vars | # classes | CART | Bagged CART | Random Forests |
|---------------|---------|--------|-----------|------|----------------|-------------------|
| Waveform | 300 | 21 | 3 | 29.1 | 19.3 | 17.2 |
| Breast Cancer | 699 | 9 | 2 | 5.9 | 3.7 | 2.9 |
| Ionosphere | 351 | 34 | 2 | 11.2 | 7.9 | 7.1 |
| Diabetes | 768 | 8 | 2 | 25.3 | 23.9 | 24.2 |
| Glass | 214 | 9 | 6 | 30.4 | 23.6 | 20.6 |

Leo Breiman (2001) “Random Forests”, Machine Learning, 45, 5-32.

Random Forests

Advantages

- Applicable to both regression and classification problems. Yes
- Handle categorical predictors naturally. Yes
- Computationally simple and quick to fit, even for large problems. Yes
- No formal distributional assumptions (non-parametric). Yes
- Can handle highly non-linear interactions and classification boundaries. Yes
- Automatic variable selection. Yes
- Very easy to interpret if the tree is small. No

Random Forests


Improve on CART with respect to:

- *Accuracy* – Random Forests is competitive with the best known machine learning methods
- *Instability* – if we change the data a little, the individual trees may change but the forest is relatively stable because it is a combination of many trees.

Random Forests and the Kinect




Random Forests and the Kinect



Random Forests and the Kinect

Use computer graphics to generate plenty of data



Shotton, et. al., Real-Time Human Pose Recognition in Parts from a Single Depth Image, CVPR 2011

Take-home messages

- **CART:** Binary decision trees can be used for **classification and regression**
- Complexity can be set by minimum samples per leaf
- **Variability in the trees:**
 - Bootstrap
 - Randomized splits
- **Averaging over multiple trees** reduces variance while bias is unaffected!



Self-test questions

You should know now:

- What we mean with non-parametric / instance-based machine learning algorithms ?
- How k-NN works ?
- How to choose the k?
- Why is it hard to use for high-D data ?
- How do search for nearest neighbours efficiently ?
- What a *binary* regression / decision tree is
- What are useful splitting criteria
- How can we influence the model complexity of the tree?
- Why is it useful to use multiple trees and randomization?