

# Chapter 2 - Unsuperised Learning

## Dimensionality Reduction and Clustering

Maschinelles Lernen 1 -  
Grundverfahren WS19/20

Prof. Gerhard Neumann  
KIT, Institut für Anthropomatik und Robotik

# Wrap-Up for Chapter 1: “Simple” Supervised Learning


| Algorithm           | Reg / Class | Representation              | Optimization / Loss                          | Pros / Cons  |
|---------------------|-------------|-----------------------------|--|--|
| Linear Regression   | Reg         | Feature Space               | Squared Error / Least Square Solution        | + Very fast, data-efficient<br>- Heavily depends on features                 |
| Logistic Regression | Class       | Feature Space               | Bernoulli log-likelihood<br>gradient Descent | + data-efficient, fast<br>- , — “ —  |
| K-NN                | Reg / Class | Instances                   | Find nearest neighbors                       | + very flexible<br>+ slow inference, curse of dim!                           |
| Random Forest       | Reg / Class | Tree / Many Trees<br>Forest | Splitting Criterion                          | + very flexible, fast<br>+ high-dimensions<br>- don't reach SOTA performance |




# Unsupervised Learning

Trainings data does not include target values, find “structure” in the data

(1) **Dimensionality reduction:**



(2) **Clustering**



(3) **Density estimation:** Generative model of the data

# Dimensionality Reduction

# Learning Outcomes

- Understand what dimensionality reduction means and why do use it
- Understand what we mean with a “projection” of a vector
- What makes a dimensionality reduction a “good” reduction
- What are the principal components in the data and what is the relation to the covariance matrix
- Learn about constraint convex optimization

# Today's Agenda!

## **Dimensionality Reduction:**

- Linear Dimensionality Reduction
- Linear Orthogonal Projections
- Reproduction Error
- Principal Component Analysis

## **Basics: Convex Constraint Optimization**

- Lagrangian Multipliers and Constraint Optimization
- Dual Optimization Problem

Slides are largely based on  
Slides from Jan Peters

# Dimensionality Reduction

## Supervised Learning:


- Learn a mapping from input  $x$  to output  $y$

## Sometimes, it is quite helpful to analyze the data points themselves

- Unsupervised learning
- Particularly:
  - Reduce the dimensionality of the data

## Possible application:

- Visualization of the data
- Preprocessing for any learning algorithm





# Motivation from Linear Least-squares Regression

- In least-squares linear regression the parameters are computed as

$$\mathbf{w} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$


where  $\mathbf{X} \in \mathbb{R}^{N \times d}$  and  $\mathbf{y} \in \mathbb{R}^{n \times 1}$

- We need to invert a  $d \times d$  matrix, which naively costs  $O(d^3)$
- Hence, it would be helpful to find a new  $d_{\text{new}} \ll d$  to gain computational advantage while not losing prediction performance

# Dimensionality Reduction

- How can we find more efficient representations for our data?
- How can we capture the “essence” of the data?


**Example: images of the digit 3**



- The images can be represented as points in a high-dimensional space (e.g., with one dimension per pixel, in a 4k image there are around 9 million dimensions!)


# Linear Dimensionality Reduction

To make things easier, we will once again assume linear models. A data point (here: one image) **can be written as a linear combination of bases** (here: basis images)



# Linear Dimensionality Reduction

- What linear transformations of the data can be used to define a lower-dimensional subspace that captures most of the structure?



# Linear Dimensionality Reduction

## Problem definition:

- Original data point  $i$ :  $\mathbf{x}_i \in \mathbb{R}^D$
- Low-dimensional representation of data point  $i$ :  $\mathbf{z}_i \in \mathbb{R}^M$  with  $D \gg M$
- **Goal:** find a mapping

$$\mathbf{x}_i \rightarrow \mathbf{z}_i$$

- Restrict this mapping to be a linear function

$$\mathbf{z}_i = \mathbf{W} \mathbf{x}_i, \text{ with } \mathbf{W} \in \mathbb{R}^{M \times D}$$

# Orthonormal Basis Vectors

We can always write a vector in terms of an orthonormal basis coordinate system

$$\mathbf{x} = \sum_{i=1}^D z_i \mathbf{u}_i, \text{ where } \mathbf{u}_i^T \mathbf{u}_j = \delta_{ij} \text{ and } \delta_{ij} = 1 \text{ if } i = j, 0 \text{ otherwise}$$

- **Orthonormality condition:** The product of 2 different basis vectors is 0. The norm of each basis vector is 1.

Example:

$$\begin{bmatrix} 3 \\ 7 \end{bmatrix} = \textcircled{3} \begin{bmatrix} 1 \\ 0 \end{bmatrix} + \textcircled{7} \begin{bmatrix} 0 \\ 1 \end{bmatrix}$$

# Projections


The coefficients  $z_i$  can be obtained by projecting  $\mathbf{x}$  on the basis vector  $\mathbf{u}_i$

$$\underbrace{z_i}_{\text{scalar coefficient}} = \underbrace{\mathbf{u}_i^T \mathbf{x}}_{\text{projection}}$$

**Example:**

$$\begin{aligned}\mathbf{x} &= z_1 \mathbf{u}_1 + z_2 \mathbf{u}_2 \\ \mathbf{u}_1^T \mathbf{x} &= z_1 \underbrace{\mathbf{u}_1^T \mathbf{u}_1}_{=1} + z_2 \underbrace{\mathbf{u}_2^T \mathbf{u}_1}_{=0} = z_1\end{aligned}$$

**Projection of 2 vectors**



# Decomposition

Use  $M \ll D$  basis vectors:

$$\mathbf{x} = \underbrace{\sum_{i=1}^M z_i \mathbf{u}_i}_{\tilde{\mathbf{x}} \approx \mathbf{x}} + \underbrace{\sum_{j=M+1}^D z_j \mathbf{u}_j}_{\text{skip}}$$

Find the  $M$  basis vectors  $\mathbf{u}_i$  that minimize the **mean squared reproduction error**:

$$\arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_M} E(\mathbf{u}_1, \dots, \mathbf{u}_M) = \arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_M} \sum_{i=1}^N \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|^2$$



# Minimizing the error

Assuming a single basis vector, the error can be written as

$$\begin{aligned} E(\mathbf{u}_1) &= \sum_{i=1}^N \|\mathbf{x}_i - \tilde{\mathbf{x}}_i\|^2 = \sum_{i=1}^N \|\mathbf{x}_i - \underbrace{(\mathbf{u}_1^T \mathbf{x}_i)}_{z_{i1}} \mathbf{u}_1\|^2 \\ &= \sum_{i=1}^N \mathbf{x}_i^T \mathbf{x}_i - 2(\mathbf{u}_1^T \mathbf{x}_i)^2 + (\mathbf{u}_1^T \mathbf{x}_i)^2 \underbrace{\mathbf{u}_1^T \mathbf{u}_1}_1 = \sum_{i=1}^N \mathbf{x}_i^T \mathbf{x}_i - \underbrace{(\mathbf{u}_1^T \mathbf{x}_i)^2}_{z_{i1}^2} \\ &= \sum_{i=1}^N \mathbf{x}_i^T \mathbf{x}_i - z_{i1}^2 \end{aligned}$$

*Handwritten notes:* A red arrow points from the term  $z_{i1}$  in the first equation to the term  $z_{i1} \mathbf{u}_1$  in the second equation. The term  $z_{i1}$  is circled in red in the first equation. The term  $\mathbf{u}_1^T \mathbf{u}_1$  is underlined in red in the second equation, with a red '1' written below it. The term  $(\mathbf{u}_1^T \mathbf{x}_i)^2$  is underlined in red in the third equation, with  $z_{i1}^2$  written below it.

# Minimizing the error

The error can be written as

$$E(\mathbf{u}_1) = \sum_{i=1}^N \mathbf{x}_i^T \mathbf{x}_i - z_{i1}^2$$

$$\Rightarrow \arg \min_{\mathbf{u}_1} E(\mathbf{u}_1) = \arg \max_{\mathbf{u}_1} \sum_{i=1}^N z_{i1}^2 = \arg \max_{\mathbf{u}_1} \sum_{i=1}^N (\mathbf{u}_1^T \mathbf{x}_i)^2$$


- Minimizing the error is equivalent to maximizing the variance of the projection. (Assuming a zero mean on the data)
- We can ensure a zero mean projection by subtracting the mean from the data

$$\bar{\mathbf{x}}_i = \mathbf{x}_i - \boldsymbol{\mu}$$

# Illustration


$$\tilde{\mathbf{x}} = \sum_{i=1}^M z_i \mathbf{u}_i + \boldsymbol{\mu}$$

- Projecting onto  $\mathbf{u}_1$  captures the **majority of the variance** and hence projecting onto it minimizes the error
- Note that these axes are **orthogonal and decorrelate** the data
  - i.e. in the coordinate frame of these axes, the data is uncorrelated (side note: this only works for Gaussians)



# Principle component analysis (PCA)

**Goal:** find the so-called **principal directions**, and the **variance** of the data along each principal direction



- $\lambda_i$  is the **marginal variance** along the **principal direction**  $\mathbf{u}_i$

# Principle component analysis

- The **first principal direction**  $\mathbf{u}_1$  is the direction along which the **variance of the projected data is maximal**

$$\mathbf{u}_1 = \arg \max_{\mathbf{u}} \frac{1}{N} \sum_{i=1}^N \left( \overbrace{\mathbf{u}^T (\mathbf{x}_i - \boldsymbol{\mu})}^{\mathbf{z}_i} \right)^2 \quad \text{s.t. } \mathbf{u}^T \mathbf{u} = 1$$

- The directions all have unit norm
- The **second principal direction** maximizes the variance of the data in the **orthogonal complement** of the first principal direction

# Derivation...

- Objective in matrix form...

$$\begin{aligned} E(u) &= \frac{1}{N} \sum_{i=1}^N \left( \underbrace{u^T}_{\text{scalar}} \underbrace{(x_i - \mu)}_{\text{scalar}} \right)^2 && Q^2 = Q \times Q^T \\ &= \frac{1}{N} \sum_{i=1}^N \left( \underbrace{u^T (x_i - \mu)}_{\text{scalar}} \underbrace{(x_i - \mu)^T u}_{\text{scalar}} \right) \\ &= u^T \underbrace{\left( \frac{1}{N} \sum_{i=1}^N (x_i - \mu)(x_i - \mu)^T \right)}_{\text{outer product} \rightarrow D \times D} u = u^T \Sigma u && \text{Matrix Form} \end{aligned}$$

- The objective can be written in terms of the **sample covariance**!

## Derivation...

We obtain the following **constrained optimization** problem

$$\mathbf{u}_1 = \arg \max_{\mathbf{u}} \mathbf{u}^T \Sigma \mathbf{u} \quad \text{s.t.} \quad \mathbf{u}^T \mathbf{u} = 1$$

We need to look at **constraint optimization** first!

# Constraint Optimization




# Basics: Constrained Optimization

Simple constrained optimization problem:  $\arg \min_x x^2$  s.t.  $x \geq b$


*objective*      *constraint*

No Constraint




$x^*=0$

$x \geq -1$



$x^*=0$

$x \geq 1$



$x^*=1$

How do we solve the constrained optimization problem? **Lagrangian Multipliers!**

# Basics: Lagrangian Multipliers

$$\min_x x^2 \quad \text{s.t. } x \geq b$$

## The Lagrangian:

- $L = \text{objective} - \text{multiplier} * \text{constraint}$

$$L(x, \lambda) = \underbrace{x^2}_{\text{objective}} - \underbrace{\lambda}_{\text{multiplier}} \cdot \underbrace{(x - b)}_{\text{constraint}}$$

## Lagrangian optimization:

$$\min_x \max_{\lambda} L(x, \lambda), \quad \text{s.t. } \lambda \geq 0$$

## Why is this equivalent?

**Min fights max!**

- $x < b$  :
  - $(x - b) < 0 \rightarrow \max_{\lambda} -\lambda(x - b) = \infty$
  - *min* won't let that happen
- $x > b$  :
  - $(x - b) > 0, \lambda \geq 0 \rightarrow \lambda^* = 0$
  - $L$  is the same as original objective
- $x = b$  :
  - $\lambda$  can be anything
  - $L$  is the same as original objective

**Min forces max to behave such that constraints are satisfied**

# General formulation

**General Formulation:**  $\min_{\mathbf{x}} f(\mathbf{x}),$   
s.t.  $h_i(\mathbf{x}) \geq b_i, \text{ for } i = 1 \dots K$

- Several inequality constraints (equality constraints also possible)

**Lagrangian optimization:**  $\min_{\mathbf{x}} \max_{\boldsymbol{\lambda}} L(\mathbf{x}, \boldsymbol{\lambda}), \quad L(\mathbf{x}, \boldsymbol{\lambda}) = f(\mathbf{x}) - \sum_{i=1}^K \lambda_i (h_i(\mathbf{x}) - b_i)$   
s.t.  $\lambda_i \geq 0, \text{ for } i = 1 \dots K$

# Dual formulation

## Primal optimization problem:

$$\begin{aligned} \min_{\mathbf{x}} f(\mathbf{x}), \\ \text{s.t. } h_i(\mathbf{x}) \geq b_i, \text{ for } i = 1 \dots K \end{aligned}$$

## Dual optimization problem:

$$\begin{aligned} \boldsymbol{\lambda}^* = \arg \max_{\boldsymbol{\lambda}} g(\boldsymbol{\lambda}), \quad g(\boldsymbol{\lambda}) = \min_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}) \\ \text{s.t. } \lambda_i \geq 0, \text{ for } i = 1 \dots K \end{aligned}$$

- $g$  is also called the **dual function** of the optimization problem
- We essentially swapped min and max in the definition of  $L$

**Slaters condition:** For a convex objective and convex constraints, solving the dual is equivalent to solving the primal!

- Optimal primal parameters can be obtained from **optimal dual parameters**, i.e.

$$\mathbf{x}^* = \arg \min_{\mathbf{x}} L(\mathbf{x}, \boldsymbol{\lambda}^*)$$

## Example:

Convex constraints:

Set of allowed solutions is convex

Convex Set

Non convex



$$\min_x x^2 \quad \text{s.t. } x \geq 1$$

1) Lagrangian  $L(x, \lambda) = x^2 - \lambda(x - 1)$

2) Find  $x^*$  - argmin  $L(x, \lambda)$ :  $\frac{\partial L}{\partial x} = 2x - \lambda \stackrel{!}{=} 0 \Rightarrow x^* = \frac{\lambda}{2}$

3) Dual-Function  $g(\lambda) = L(x^*, \lambda) = \frac{\lambda^2}{4} - \lambda\left(\frac{\lambda}{2} - 1\right) = -\frac{\lambda^2}{4} + \lambda$

4) Find  $\lambda^*$  - argmax  $g(\lambda)$ :  $\frac{\partial g}{\partial \lambda} = -\frac{\lambda}{2} + 1 \stackrel{!}{=} 0 \Rightarrow \lambda^* = 2$

5)  $x^* = \frac{\lambda^*}{2} = 1$



# Back to the PCA Derivation...

We obtain the following **constrained optimization** problem

$$\mathbf{u}_1 = \arg \max_{\mathbf{u}} \mathbf{u}^T \Sigma \mathbf{u} \quad \text{s.t.} \quad \mathbf{u}^T \mathbf{u} = 1$$

- We now know what to do... **Lagrangian optimization**

The **Lagrangian** is given by:

$$L(\mathbf{u}, \lambda) = \mathbf{u}^T \Sigma \mathbf{u} - \lambda(\mathbf{u}^T \mathbf{u} - 1)$$

- Optimal solution for  $\mathbf{u}$ :

$$\frac{\partial L(\mathbf{u}, \lambda)}{\partial \mathbf{u}} = \cancel{2}\Sigma \mathbf{u} - \cancel{2}\lambda \mathbf{u} \stackrel{!}{=} \mathbf{0} \quad \Rightarrow \quad \Sigma \mathbf{u} = \lambda \mathbf{u} \quad \text{This is an **Eigen-value problem!**}$$

# Basics: Eigenvalues and Eigenvectors

- Let the **Eigenvectors** and **Eigenvalues** of  $\mathbf{C}$  be  $\mathbf{u}_k$  and  $\lambda_k$  for  $k \leq D$  i.e.,

$$\mathbf{C}\mathbf{u}_k = \lambda_k \mathbf{u}_k \quad \text{with } \lambda_1 \geq \lambda_2 \geq \dots \geq \lambda_D \quad \text{Ordered list of Eigenvalues}$$

- In matrix form:

$$\mathbf{C}\mathbf{U} = \mathbf{U}\mathbf{\Lambda} \quad \text{with } \mathbf{\Lambda} = \text{diag}(\lambda_1, \dots, \lambda_D) \text{ and } \mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_D]$$

- Because  $\mathbf{U}$  is **orthonormal** (eigenvectors have unit norm), we know that  $\mathbf{U}\mathbf{U}^T = \mathbf{I}$
- This mean that we can decompose  $\mathbf{C}$  as

$$\underbrace{(\mathbf{C}\mathbf{U})}_{\mathbf{I}} \mathbf{U}^T = (\mathbf{U}\mathbf{\Lambda})\mathbf{U}^T \Rightarrow \mathbf{C} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T$$



# Basics: Eigenvalues and Eigenvectors

Every **positive definite symmetric matrix** can be decomposed in its **Eigendecomposition**

$$C = U\Lambda U^T = \underbrace{\begin{bmatrix} \mathbf{u}_1 & \dots & \mathbf{u}_D \end{bmatrix}}_{\text{Eigenvectors}} \underbrace{\begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_D \end{bmatrix}}_{\text{Eigenvalues}} \begin{bmatrix} \mathbf{u}_1^T \\ \vdots \\ \mathbf{u}_D^T \end{bmatrix}$$

# Back to PCA

## Eigenvalues-Eigenvectors of the covariance matrix

$$\Sigma u = \lambda u$$

- The **largest** Eigenvalue gives us the **maximal variance**
- The corresponding Eigenvector gives us the **direction with maximal variance**

# Principal Component Analysis

- **Observation:** If  $\lambda_k \approx 0$  for  $k > M$  for some  $M \ll D$ , then we can use the subset of the first  $D$  eigenvectors to define a basis for approximating the data vectors with loosing accuracy

$$\mathbf{x}_i - \boldsymbol{\mu} = \underbrace{\sum_{j=1}^M z_{ij} \mathbf{u}_j}_{\tilde{\mathbf{x}}} + \underbrace{\sum_{j=M+1}^D z_{ij} \mathbf{u}_j}_{\text{close to 0}} \Rightarrow \mathbf{x}_i \approx \boldsymbol{\mu} + \sum_{j=1}^M z_{ij} \mathbf{u}_j$$

- This representation has the **minimal mean squared error** (MSE) of all linear representations of dimension  $D$

$$\arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_M} E(\mathbf{u}_1, \dots, \mathbf{u}_M) = \arg \min_{\mathbf{u}_1, \dots, \mathbf{u}_M} \sum_{i=1}^N ||\mathbf{x}_i - \tilde{\mathbf{x}}_i||^2$$

# Principal Component Analysis

Now we know how we can represent our data in a **lower dimensional space** in a **principled way**

- **Center** the data around the mean (compute the mean of the data and subtract it)
- Compute the **covariance matrix**, decompose it, and choose the **first M largest** Eigenvalues and corresponding Eigenvectors
- This gives us an **(Eigen)basis** for representing the data

- **Projection to low-D:**  $\mathbf{z}_i = \mathbf{B}^T (\mathbf{x}_i - \boldsymbol{\mu})$

- **Reprojection to high-D:**  $\tilde{\mathbf{x}}_i = \boldsymbol{\mu} + \mathbf{B} \mathbf{z}_i$


with  $\mathbf{B} = \begin{bmatrix} \mathbf{u}_1 & \dots & \mathbf{u}_M \end{bmatrix}$

- It is also common to normalize the variance of each dimension (i.e. unit variance)


# How to choose M

- A larger M leads to a better approximation. In the limit, when  $M = D$  we stay in the initial data dimensions
- There are at least 2 good possibilities for choosing M
  - Choose D based on **application performance**, i.e. choose the smallest M that makes the application work well enough
  - Choose D so that the **Eigenbasis captures some fraction of the variance** (for example  $\eta = 0.9$ ).  
The eigenvalue  $\lambda_i$  describes the marginal variance captured by  $\mathbf{u}_i$


$$\text{Choose } D \text{ s.t. } \sum_{i=1}^M \lambda_i = \eta \underbrace{\sum_{i=1}^D \lambda_i}_{\text{Total variance of the data}}$$



# Image representation with PCA



# Image representation with PCA




$$= \text{[Image 1]} + a_1 \text{[Image 2]} + a_2 \text{[Image 3]} + a_3 \text{[Image 4]}$$

The equation shows the reconstruction of an image as a linear combination of principal components. The first image is the mean image (green border). The subsequent images are the principal components (red, purple, and blue borders). The coefficients  $a_1$ ,  $a_2$ , and  $a_3$  are shown in red, purple, and teal respectively.


# Eigenfaces

- The first popular use of PCA for object recognition was for the detection and recognition of faces [Turk and Pentland, 1991]
- Collect a face ensemble
- Normalize for contrast, scale, & orientation
- Remove backgrounds
- Apply PCA & choose the first M eigen-images that account for most of the variance of the data





# Image Morphing with PCA




# Generic Image Ensembles

Is there a low-dimensional model describing natural images?




# PCA of natural image patches

8x8 image patches



# PCA Model of body shapes

- PCA on a detailed triangle model of human bodies [Anguelov et al. 05]



# Wrap-up

## Summary:

- PCA projects the data into a linear subspace
- PCA maximizes the variance of the projection
- PCA minimizes the error of the reconstruction
- We just covered the most simple linear dimensionality reduction technique
  - Many more sophisticated techniques exist
  - Kernel PCA, Auto-Encoders, t-SNE, non-negative matrix factorization (interesting, but no time to cover those...)

## Applications:

- PCA allows us to transform a high-dimensional input space to a low-dimensional feature space, while capturing the essence of the data
- PCA finds a more natural coordinate system for the data
- PCA is a **very common preprocessing step** for high-dimensional input data

# Self-test questions

## **What have we learned today?**



- What does dimensionality reduction mean?
- How does linear dimensionality reduction work?
- What is PCA? What are the three things that it does?
- What are the roles of the Eigenvectors and Eigenvalues in PCA?
- Can you describe applications of PCA?

# Clustering

Slides based on lectures from Stefan Roth (TU Darmstadt)  
and Carla Brodley (Tulft University)

# Clustering is Subjective


What is the natural grouping of these objects?




Simpson's Family



School Employees



Females



Males




# Clustering: Finding structure in the data

## What are the correct clusters?

- Ground truth often not available

## Similarity measure

- clustering relies on measure of similarity
- e.g. position in space (Euclidean vs. log-polar coordinates), weighting of different dimensions (features)...




# What is similarity?





Similarity is hard to  
define, but...  
*"We know it when we  
see it"*

# Defining Distance Measures

- Definition:** Let  $O_1$  and  $O_2$  be two objects from the universe of possible objects. The distance (dissimilarity) between  $O_1$  and  $O_2$  is a real number denoted by  $D(O_1, O_2)$



Peter Piotr




# What properties should a distance measure have?

- $D(A,B) = D(B,A)$  **Symmetry**  
*Otherwise you could claim “Alex looks like Bob, but Bob looks nothing like Alex.”*
- $D(A,A) = 0$  **Constancy of Self-Similarity**  
*Otherwise you could claim “Alex looks more like Bob, than Bob does.”*
- $D(A,B) = 0$  iif  $A = B$  **Positivity (Separation)**  
*Otherwise there are objects in your world that are different, but you cannot tell apart.*
- $D(A,B) \leq D(A,C) + D(B,C)$  **Triangular Inequality**  
*Otherwise you could claim “Alex is very like Bob, and Alex is very like Carl, but Bob is very unlike Carl.”*

# Basic Clustering Algorithms

## Hierarchical clustering methods

- Bottom-up (merging)
- Top-down (splitting, not covered)




---

## Flat clustering algorithms

- K-Means
- Mixture models (see density estimation lecture)

## Other clustering methods:


- Spectral clustering (not covered)




# Hierarchical Clustering: Dendograms

## Dendogram:


- A useful tool for summarizing similarity measurements




- The similarity between two objects in a dendrogram is represented as the **height of the lowest internal node they share**



# Example: Species based on genetic difference




## Example: Names based on string edit distance






# Properties of Dentograms

We can look at the dendrogram to determine the “correct” number of clusters.



# Properties of Dentograms

**Detecting Outliers:** The single isolated branch is suggestive of a data point that is very different to all others



# Hierarchical Clustering

Since we cannot test all possible trees we will have to heuristic search of all possible trees. We could do this..

- **Bottom-Up (agglomerative):** Starting with each item in its own cluster, find the best pair to merge into a new cluster. Repeat until all clusters are fused together.
- **Top-Down (divisive):** Starting with all the data in a single cluster, consider every possible way to divide the cluster into two. Choose the best division and recursively operate on both sides. (not covered)

# Bottom-Up Clustering

## Hierarchical agglomerative clustering

- sequentially merge the pairs of closest points

### Init:


- Each of the  $n$  samples becomes a cluster itself

### Iterate:

- (1) find closest clusters and merge them
- (2) proceed until we have a single cluster

### Required:

- Distance measure between two samples
- Distance measure between clusters



# Similarity between clusters



**Cluster Linkage:** define distances between clusters

- **Single Linkage:** Minimum distance between two points from the 2 cluster

$$d(C_k, C_l) = \min_{\mathbf{x}_i \in C_k} \min_{\mathbf{x}_j \in C_l} d(\mathbf{x}_i, \mathbf{x}_j)$$

- **Complete Linkage:** Maximum distance between two points from the 2 cluster

$$d(C_k, C_l) = \max_{\mathbf{x}_i \in C_k} \max_{\mathbf{x}_j \in C_l} d(\mathbf{x}_i, \mathbf{x}_j)$$



# Similarity between clusters



**Cluster Linkage:** define distances between clusters

- **Average Linkage:** Average distance between any 2 pairs



$$d(C_k, C_l) = \frac{1}{|C_l||C_k|} \sum_{\mathbf{x}_i \in C_l} \sum_{\mathbf{x}_j \in C_k} d(\mathbf{x}_i, \mathbf{x}_j)$$

- **Centroid Linkage:** Distance between the 2 centroids


$$d(C_k, C_l) = d\left(\frac{1}{|C_l|} \sum_{\mathbf{x}_i \in C_l} \mathbf{x}_i, \frac{1}{|C_k|} \sum_{\mathbf{x}_j \in C_k} \mathbf{x}_j\right)$$



# Influence of linkage method



Single linkage



Average linkage

# Hierarchical Clustering Methods Summary

## Wrap-Up:


- No need to specify the number of clusters in advance
- Hierarchical nature maps nicely onto human intuition for some domains
- They do not scale well: time complexity of at least  $O(n^2)$ , where  $n$  is the number of total objects
- Like any heuristic search algorithms, local optima are a problem
- Interpretation of results is (very) subjective



# Basic Clustering Algorithms

## Hierarchical clustering methods

- Bottom-up (merging)
- Top-down (splitting, not covered)




---

## Flat clustering algorithms

- K-Means
- Mixture models (see density estimation lecture)

## Other clustering methods:

- Spectral clustering (not covered)




# K-Means Algorithm

## Goal: minimize quantization error!

- Given data  $D = \{\mathbf{x}_1, \dots, \mathbf{x}_n\}$
- Search for **cluster centers/prototypes/centroid**  
 $C = \{\mathbf{c}_1, \dots, \mathbf{c}_k\}$
- Denote with  $c(\mathbf{x})$  the closest centroid vector  $\mathbf{c} \in C$  to  $\mathbf{x}$
- Here sum of squared distances (SSD) (or sum of squared error) denotes quantization error

$$\text{SSD}(C; \mathcal{D}) = \sum_{i=1}^n d(\mathbf{x}_i, c(\mathbf{x}_i))^2$$



# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step 1:** The stars are cluster centers, randomly assigned at first.

# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step 2:** Assign each example to its nearest cluster center.

# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step 3:** Adjust the **centroids to be the means** of the samples assigned to them.

# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step 4:** Assign each example to its nearest cluster center.

# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step 5:** Adjust the **centroids to be the means** of the samples assigned to them.

# K-Means Algorithm

## Iterative Procedure


1. Pick K arbitrary cluster centers
2. Assign each sample to its closest centroid

$$z_n = \arg \min_k ||\mathbf{c}_k - \mathbf{x}_n||^2$$

3. Adjust the **centroids to be the means** of the samples assigned to them

$$\mathbf{c}_k = \frac{1}{|X_k|} \sum_{\mathbf{x}_i \in X_k} \mathbf{x}_i, \quad X_k = \{\mathbf{x}_n | z_n = k\}$$

4. Goto step 2 until no change



**Step N: Convergence...**



# Does K-Means converge?

To analyze convergence, we write SSD in terms of assignments  $z_n$

$$\text{SSD}(C; \mathcal{D}) = \sum_{i=1}^n d(\mathbf{x}_i, c(\mathbf{x}_i))^2 = \sum_{i=1}^n \sum_k q_{nk} d(\mathbf{x}_i, \mathbf{c}_k)^2,$$

where  $q_{nk} = \mathbb{I}(z_n = k)$  is 1 if the  $n$ -th example is assigned to the  $k$ -th cluster and 0 otherwise (1-hot coding)

- **Assignment Step:** Minimizes SSD w.r.t.  $z_n$ 
  - Sets  $q_{nk}$  of nearest cluster to 1, all other values are 0
- **Adjustment Step:** Minimizes SSD w.r.t. centroids  $\mathbf{c}_k$

$$\mathbf{c}_k = \frac{1}{\sum_n q_{nk}} \sum_{n=1}^N q_{nk} \mathbf{x}_n$$

- Average vector is the vector with minimum squared distance to all assigned samples

# K-Means analysis

**Does K-Means converge?** Yes, it (locally) minimizes the SSD!

- We have only a finite number of possible values for the centroid
- Every assignment or adjustment step is reducing the SSD (or it stays constant)

**Does K-Means converge to the global minimal cost solution?** No!

- The objective is an NP-Hard problem, so we can't expect any algorithm to minimize the cost without essentially checking (near to) all assignments.
- It heavily depends on the initialisation of the centroids


# K-means++

## Furthest First Initialization:

- Pick a random data-point as first center
- **for**  $k \in \{2, \dots, K\}$  **do**
  - find the example that is furthest from all previously selected means

$$\text{let } n = \arg \max_{n \in \{1, \dots, N\}} \left( \min_{k' \in \{1, \dots, k-1\}} \|x_n - c_{k'}'\|^2 \right)$$

- Assign centroid:  $c_k = x_n$



## Furthest first initialization in action...


- Converges (in this case) after 1 adjustment

# Number of clusters


## How to choose K?

- Based on 'good' function value decrease on 'holdout' set, cross validation (good but expensive)
- "Knee-finding method" (similar to PCA, heuristic but cheap)


When  $k = 1$ , the objective function is 873.0



When  $k = 2$ , the objective function is 173.1




When  $k = 3$ , the objective function is 133.6



## “Knee-finding” method

- We can plot the objective function (SSD) values for  $k$  equals 1 to 6...
- SSD will decrease with higher  $k$  (on average)
- The abrupt change at  $k = 2$ , is highly suggestive of two clusters in the data.
- This technique for determining the number of clusters is known as “knee finding” or “elbow finding”



# Wrap-Up

## **Strengths:**

- K-means usually converges very quickly in practice.
- K-means++ still not guaranteed to find the global optima
  - in practice, we can get stuck.
  - often try multiple initializations (use a little randomness in K-means++ and run the algorithm multiple times).

## **Weaknesses:**

- Applicable only when mean is defined, then what about categorical data?
- Unable to handle noisy data and outliers
- Not suitable to discover clusters with non-convex shapes

# Self-test Questions

## **What you should know now:**

- How is the clustering problem defined? Why is it called “unsupervised”?
- How do hierarchical clustering methods work? What is the rule of the cluster-2-cluster distance and which distances can we use?
- How does the k-mean algorithm work? What are the 2 main steps?
- Why does the algorithm converge? What is it minimizing?
- Does k-means finds a the global minimum of the objective?